

POWE.RS Mathematical Formulations

Document Purpose: Complete mathematical specification of the SDDP algorithm, LP subproblem formulation, stochastic modeling, and convergence analysis for the POWE.RS hydrothermal dispatch solver.

Notation Convention: Follows [SDDP.jl](#) canonical notation.

Last Updated: 2026-01-22

Table of Contents

Part I: SDDP Algorithm Foundation

1. Introduction
2. SDDP Algorithm Overview

Part II: Stage Subproblem Formulation

3. Notation and Sets
4. Base LP Formulation
5. Block Formulation Variants
6. Hydro Production Function Models
7. Equipment-Specific Formulations

Part III: Stochastic Modeling

8. PAR(p) Inflow Model
9. Inflow Non-Negativity Methods

Part IV: Cut Management and Convergence

10. Cut Generation and Aggregation
11. Cut Selection Strategies
12. Stopping Rules

Part V: Advanced Formulations

13. Discount Rate
14. Infinite Periodic Horizon
15. Upper Bound Evaluation (Inner Approximation)
16. Risk-Averse SDDP (CVaR)

Part VI: Configuration Reference

17. Configuration-Driven LP Variants
18. Cross-Reference to Data Model

Appendices

- [Appendix A: Notation Reference](#)
 - [Appendix B: SDDP Algorithm Pseudocode](#)
 - [Appendix C: Deferred Features](#)
-

Part I: SDDP Algorithm Foundation

1. Introduction

1.1 Document Purpose

This document provides the complete mathematical specification for the POWE.RS Stochastic Dual Dynamic Programming (SDDP) solver. It serves as the authoritative reference for:

- **Algorithm specification:** The SDDP iterative process, including forward and backward passes
- **LP formulation:** The stage subproblem structure, decision variables, constraints, and objective
- **Stochastic modeling:** Inflow models, scenario trees, and uncertainty handling
- **Convergence analysis:** Stopping rules, bound computation, and gap analysis
- **Extensions:** Risk measures, infinite horizon, and advanced features

For data structures, file formats, and implementation details, see [DATA_MODEL_SPECIFICATION.md](#).

1.2 Notation Conventions

This document follows [SDDP.jl](#) notation conventions for consistency with the broader SDDP literature:

Convention	Meaning
$t \in \{1, \dots, T\}$	Stage index
$\omega \in \Omega_t$	Scenario realization at stage t
x_t	State variables at end of stage t
$\hat{x}_{t-1}$	Incoming state (from previous stage)
$V_t(x)$	Value function (cost-to-go) at stage t
θ_t	Epigraph variable approximating V_{t+1}
π	Dual variables (Lagrange multipliers)
(α, β)	Cut intercept and coefficients

1.3 Problem Context

POWE.RS solves the **hydrothermal dispatch problem**: determining optimal generation schedules for hydro and thermal plants over a multi-year planning horizon under inflow uncertainty. Key characteristics:

- **Long horizons:** 5-10 years (60-120 monthly stages)
- **Large state space:** 160+ hydro reservoirs with AR inflow models
- **Stochastic inflows:** PAR(p) autoregressive models with seasonal patterns
- **Coupling constraints:** Cascade hydrology, transmission limits, energy balance

2. SDDP Algorithm Overview

2.1 Multistage Stochastic Programming Formulation

The hydrothermal dispatch problem is formulated as a multistage stochastic program:

$$\min_{x_1, \dots, x_T} \mathbb{E} \left[\sum_{t=1}^T c_t(\omega_t)^\top x_t(\omega_{[t]}) \right]$$

subject to stage-linking constraints and uncertainty realization. The nested formulation uses **value functions**:

$$V_t(x_{t-1}) = \mathbb{E}_{\omega_t} \left[\min_{x_t} \{ c_t^\top x_t + V_{t+1}(x_t) : A_t x_t = b_t - E_t x_{t-1}, x_t \in \mathcal{X}_t \} \right]$$

with terminal condition $V_{T+1}(x) = 0$.

Key insight: The value function $V_t(x)$ is convex and piecewise linear (for LP subproblems), enabling outer approximation via Benders cuts.

2.2 The SDDP Algorithm

SDDP iteratively builds piecewise-linear approximations $\hat{V}_t^k$ of the true value functions through:

1. **Forward pass:** Sample scenarios, make decisions using current approximation
2. **Backward pass:** Compute cuts to improve the approximation
3. **Convergence check:** Evaluate stopping criteria

2.2.1 Forward Pass The forward pass simulates the system under the current policy to generate **trial points** (visited states):

Algorithm: Forward Pass (iteration k , pass m)

- **Input:** Initial state x_0 , cut approximations $\{\hat{V}_t^k\}$
 - **Output:** Visited states $\{\hat{x}_t^m\}_{t=1}^T$, scenario costs
1. Set $\hat{x}_0 = x_0$
 2. For $t = 1$ to T :
 - Sample $\omega_t \sim P(\Omega_t)$
 - Solve stage LP with incoming state $\hat{x}_{t-1}$ and realization ω_t :

$$\hat{x}_t, \hat{\theta}_t = \arg \min \{c_t^\top x_t + \theta_t : \text{constraints}(x_t, \hat{x}_{t-1}, \omega_t), \theta_t \geq \alpha_i + \beta_i^\top x_t \ \forall \text{ cut } i\}$$

- Record visited state $\hat{x}_t$
3. Return $\{\hat{x}_t^m\}_{t=1}^T$

Parallelization: Forward passes are embarrassingly parallel—each scenario trajectory is independent. POWERS distributes M forward passes across MPI ranks.

2.2.2 Backward Pass The backward pass computes cuts by walking stages in reverse order:

Algorithm: Backward Pass (iteration k)

- **Input:** Visited states from forward passes
 - **Output:** New cuts for each stage
1. For $t = T$ down to 1:
 - For each visited state $\hat{x}_{t-1}$ from forward passes:
 - For each $\omega \in \Omega_t$ (branching scenarios):
 - * Solve stage LP with $(\hat{x}_{t-1}, \omega)$
 - * Extract: $Q_t(\hat{x}_{t-1}, \omega) = \text{optimal value}$
 - $\pi_t(\omega) = \text{dual of state constraints}$
 - * Compute per-scenario cut coefficients (see Section 3.4 for sign convention):
 - $\beta(\omega) = \pi_t(\omega)$ (for state constraints $x_t = \hat{x}_{t-1} + \dots$)
 - $\alpha(\omega) = Q_t - \beta(\omega)^\top \hat{x}_{t-1}$
 - Aggregate cut (single-cut formulation):
 - $\bar{\beta} = \sum_{\omega} p(\omega) \cdot \beta(\omega)$
 - $\bar{\alpha} = \sum_{\omega} p(\omega) \cdot \alpha(\omega)$
 - Add cut to stage $t - 1$:
 - $\theta_{t-1} \geq \bar{\alpha} + \bar{\beta}^\top x_{t-1}$

Warm-starting: The forward pass solution provides a near-optimal basis for backward branching scenarios, significantly reducing solve times.

2.2.3 Convergence Monitoring Lower Bound: The deterministic lower bound is the first-stage LP value:

$$z^k = V_1^k(x_0) = \min_{x_1} \{c_1^\top x_1 + \theta_1 : \text{constraints}, \theta_1 \geq \alpha_i + \beta_i^\top x_1 \ \forall i\}$$

This bound increases monotonically as cuts are added.

Upper Bound: Estimated via Monte Carlo simulation or inner approximation (see Section 15):

$$\bar{z}^k = \frac{1}{M} \sum_{m=1}^M \sum_{t=1}^T c_t^\top x_t^{(m)}$$

Optimality Gap:

$$\text{gap}^k = \frac{\bar{z}^k - z^k}{\max(1, |\bar{z}^k|)}$$


2.3 Policy Graph Structure

2.3.1 Finite Horizon (Acyclic Graph) The standard SDDP formulation uses an acyclic directed graph:



- **Nodes:** Stages $t \in \{1, \dots, T\}$
- **Arcs:** Transitions with probabilities (typically deterministic: $p = 1$)
- **Terminal:** $V_{T+1}(x) = 0$ (no future cost)

2.3.2 Cyclic Graph (Infinite Horizon) For long-term planning, POWE.RS supports **infinite periodic horizon** with cyclic graphs:



- **Cycle:** Stage T transitions back to stage 1 (or a cycle start)
- **Discount:** Cycle transitions require discount rate $\beta < 1$ for convergence
- **Cut sharing:** Cuts at equivalent cycle positions are shared

See Section 14 for the complete infinite horizon formulation.

2.4 State Variables and the Markov Property

For SDDP to generate valid cuts, the subproblem must satisfy the **Markov property**: future costs depend only on the current state, not on how we arrived at that state.

State variables in POWE.RS:

Component	Variable	Count	Description
Hydro storage	v_h	N_{hydro}	Reservoir volume at end of stage
AR inflow lags	$a_{h,\ell}$	$\sum_h P_h$	Lagged inflows for AR(P) models
Battery SOC	soc_{bat}	$N_{battery}$	Battery state of charge (DEFERRED)
GNL pipeline	$gnl_{t,\ell}$	$\sum_{gnl} L_{gnl}$	Committed GNL dispatch (DEFERRED)

State expansion trick: The AR inflow model requires past inflows $a_{h,t-1}, a_{h,t-2}, \dots$ to compute current inflow. To maintain the Markov property, these lags are included as state variables with trivial dynamics:

$$a_{h,\ell}^{out} = a_{h,\ell-1}^{in} \quad \text{for } \ell = 2, \dots, P_h$$

This allows cut coefficients to capture the value of inflow history.

2.5 Single-Cut vs Multi-Cut Formulation

Single-Cut (Default): One aggregated cut per iteration:

$$\theta_{t-1} \geq \bar{\alpha} + \bar{\beta}^\top x_{t-1}$$

where $\bar{\alpha} = \mathbb{E}[\alpha(\omega)]$ and $\bar{\beta} = \mathbb{E}[\beta(\omega)]$.

- **Pros:** Fewer cuts, smaller LP, faster solves
- **Cons:** May require more iterations to converge

Multi-Cut (DEFERRED): One cut per scenario per iteration:

$$\theta_{t-1,\omega} \geq \alpha(\omega) + \beta(\omega)^\top x_{t-1} \quad \forall \omega \in \Omega_t$$

- **Pros:** Tighter approximation, fewer iterations
- **Cons:** More cuts, larger LP, memory-intensive

POWE.RS implements single-cut by default. Multi-cut is planned for future implementation.

Part II: Stage Subproblem Formulation

3. Notation and Sets

3.1 Index Sets

Symbol	Description	Typical Size	Notes
$t \in \{1, \dots, T\}$	Stages	60-120	5-10 year monthly horizon
$k \in \mathcal{K}$	Blocks within stage	1-24	3 typical (LEVE/MÉDIA/PESADA)
$\mathcal{B}$	Buses	4-10	4-5 for SIN subsystems
$\mathcal{H}$	Hydro plants	160	All plants in system
$\mathcal{H}^{op} \subseteq \mathcal{H}$	Operating hydros (can generate)	$\approx \mathcal{H} $	Most/all plants typically operating
$\mathcal{H}^{fill} \subseteq \mathcal{H}$	Filling hydros (no generation)	0	Usually 0; rare for new plants under commissioning
$\mathcal{T}$	Thermal plants	130	
$\mathcal{L}$	Transmission lines	10	Regional interconnections
$\mathcal{C}^{imp}, \mathcal{C}^{exp}$	Import/export contracts	5	
$\mathcal{P}$	Pumping stations	5	
$\mathcal{G}$	Generic constraints	50	User-defined
$\mathcal{S}_b$	Deficit segments for bus b	1	Multiple segments optional

Symbol	Description	Typical Size	Notes
$\mathcal{M}_h$	FPHA planes for hydro h	125	Typical value; depends on grid resolution
$\mathcal{U}_h$	Upstream hydros of h	1-2	Immediate upstream in cascade
Ω_t	Scenario realizations at stage t	20	Standard NEWAVE branching factor

3.2 Parameters

Time and Conversion:

Symbol	Units	Description
τ_k	hours	Duration of block k
$w_k = \tau_k / \sum_j \tau_j$	-	Block weight (fraction of stage)
ζ	hm ³ /(m ³ /s)	Time conversion: m ³ /s over stage $\rightarrow$ hm ³

Time Conversion Factor Derivation The factor ζ converts a flow rate in m³/s to a volume in hm³ accumulated over the stage duration.

Fundamental Relationship:

$$\text{Volume} = \text{Flow Rate} \times \text{Time}$$

Unit Conversion Chain: 1. Flow rate: Q [m³/s] 2. Time period: τ [hours] 3. Target volume: V [hm³] = 10⁶ m³

$$V \text{ [hm}^3\text{]} = Q \text{ [m}^3\text{/s]} \times \tau \text{ [hours]} \times \frac{3600 \text{ s}}{1 \text{ hour}} \times \frac{1 \text{ hm}^3}{10^6 \text{ m}^3}$$

$$V = Q \times \tau \times \frac{3600}{10^6} = Q \times \tau \times 0.0036$$

For a stage with multiple blocks: If the stage has blocks $k \in \mathcal{K}$ with durations τ_k hours, and the flow is assumed constant across the stage (parallel blocks), the total time is $\sum_k \tau_k$ hours:

$$\zeta = 0.0036 \times \sum_{k \in \mathcal{K}} \tau_k \text{ [hm}^3\text{/(m}^3\text{/s)]}$$

Dimensional Analysis:

$$[\zeta] = \frac{\text{s}}{\text{h}} \times \frac{\text{m}^3}{\text{hm}^3} \times \text{h} = \frac{\text{hm}^3}{\text{m}^3/\text{s}}$$

Worked Example (Monthly Stage):

Block	Name	Duration τ_k (h)
1	LEVE	200
2	MÉDIA	300
3	PESADA	228
Total		728

$$\zeta = 0.0036 \times 728 = 2.6208 \text{ hm}^3/(\text{m}^3/\text{s})$$

Verification: A constant inflow of $Q = 100 \text{ m}^3/\text{s}$ over the month yields:

$$V = Q \times \zeta = 100 \times 2.6208 = 262.08 \text{ hm}^3$$

Direct calculation: $100 \text{ m}^3/\text{s} \times 728 \text{ h} \times 3600 \text{ s/h}/10^6 = 262.08 \text{ hm}^3$

Load and Costs:

Symbol	Units	Description
$D_{b,k}$	MW	Load at bus b , block k
$c_{b,s}^{def}$	\$/MWh	Deficit cost at bus b , segment s
$\bar{d}_{b,s}$	MW	Deficit segment depth
c_b^{exc}	\$/MWh	Excess generation cost
$c_{t,s}^{th}$	\$/MWh	Thermal cost at plant t , segment s
c_h^{spill}	\$/(\text{m}^3/\text{s} \cdot \text{h})	Spillage cost
c_h^{div}	\$/(\text{m}^3/\text{s} \cdot \text{h})	Diversion cost
c_l^{exch}	\$/MWh	Exchange (transmission) cost
c_j^{pump}	\$/MWh	Pumping cost
c_c^{imp}, c_c^{exp}	\$/MWh	Contract import cost / export revenue

Hydro Parameters:

Symbol	Units	Description
$\hat{v}_h$	hm^3	Incoming storage (state from previous stage)
$\bar{V}_h, \underline{V}_h$	hm^3	Storage bounds
$\bar{Q}_h, \underline{Q}_h$	m^3/s	Turbined flow bounds
$\bar{G}_h, \underline{G}_h$	MW	Generation bounds
$\bar{O}_h, \underline{O}_h$	m^3/s	Outflow bounds
ρ_h	$\text{MW}/(\text{m}^3/\text{s})$	Productivity (constant model)
$\gamma_0^m, \gamma_v^m, \gamma_q^m, \gamma_s^m$	-	FPHA plane coefficients

Transmission and Contract Parameters:

Symbol	Units	Description
$\bar{F}_l^+, \bar{F}_l^-$	MW	Line capacity (direct/reverse)
$\eta_l = 1 - \text{losses}/100$	-	Line efficiency
$\bar{M}_c$	MW	Contract capacity

Inflow Model Parameters:

Note on Periodicity: The PAR(p) model uses periodic parameters that repeat with a cycle length M . Common configurations: - **Monthly stages:** $M = 12$ (seasons = months) - **Weekly stages:** $M = 52$ (seasons = weeks) - **Custom resolution:** $M = \text{number of distinct periods in the cycle}$

We use “**season m** ” as a generic term for the position within the cycle, avoiding the term “month” which is resolution-specific. The mapping $m(t) = ((t - 1) \bmod M) + 1$ converts stage index t to season index $m \in \{1, \dots, M\}$.

Symbol	Units	Description
μ_m	m^3/s	Seasonal mean inflow for season m
$\psi_{m,\ell}$	-	AR coefficient for season m , lag ℓ
σ_m	m^3/s	Residual standard deviation for season m
$\hat{a}_{h,\ell}$	m^3/s	Incoming AR lag ℓ (state)

3.3 Decision Variables

Notation Convention: - **Generation variables** use g with entity subscript: g_h (hydro at plant h), g_j (thermal at plant j) - **Flow variables** use intuitive single letters: q (turbined), s (spillage), u (diversion/bypass) - **Total outflow** is explicitly defined: $o_h = q_h + s_h$ (downstream channel flow) - **Contract variables** use χ (chi) with direction superscript: χ^{in} , χ^{out} - **Slack variables** use σ with constraint-type superscript

Symbol Selection Rationale: - q (turbined): from “vazão turbinada” (Portuguese) or “discharge through turbines” - s (spillage): standard hydrology notation - u (diversion): “bypass” or “desvio” — avoids confusion with demand D or deficit δ - o (outflow): total downstream flow affecting tailrace level - r (withdrawal): “retirada” — consumptive removal from the system - χ (contract): Greek chi, visually distinct from cost symbol c

Per-Block Variables (indexed by $k \in \mathcal{K}$):

Variable	Domain	Units	Description
$\delta_{b,k,s}$	$[0, \bar{d}_{b,s}]$	MW	Deficit at bus b , segment s
$\epsilon_{b,k}$	≥ 0	MW	Excess generation at bus b
$f_{l,k}^+$	$[0, \bar{F}_l^+]$	MW	Direct flow on line l
$f_{l,k}^-$	$[0, \bar{F}_l^-]$	MW	Reverse flow on line l
$g_{j,k,s}$	$[0, \bar{g}_{j,s}]$	MW	Thermal generation at plant j , segment s
$q_{h,k}$	$[\underline{Q}_h, \bar{Q}_h]$	m^3/s	Turbined flow at hydro h
$s_{h,k}$	≥ 0	m^3/s	Spillage at hydro h
$g_{h,k}$	$[\underline{G}_h, \bar{G}_h]$	MW	Hydro generation at plant h
$u_{h,k}$	$[0, \bar{U}_h]$	m^3/s	Diversion/bypass flow (to separate channel)
$o_{h,k}$	-	m^3/s	Total downstream outflow: $o_{h,k} = q_{h,k} + s_{h,k}$
$e_{h,k}$	free	m^3/s	Evaporation (can be negative for condensation)
$r_{h,k}$	≥ 0	m^3/s	Water withdrawal (consumptive use)
$p_{j,k}$	$[0, \bar{P}_j]$	m^3/s	Pumped flow at station j
$\chi_{c,k}^{in}$	$[0, \bar{C}_c]$	MW	Contract import
$\chi_{c,k}^{out}$	$[0, \bar{C}_c]$	MW	Contract export

Stage-Level State Variables:

Variable	Domain	Units	Description
v_h	$[V_h, \bar{V}_h]$	hm ³	End-of-stage storage
v_h^{avg}	-	hm ³	Average storage during stage: $(\hat{v}_h + v_h)/2$
$a_{h,\ell}$	fixed	m ³ /s	AR lag ℓ (fixed by state transition)
θ	≥ 0	\$	Future cost (cost-to-go approximation)

Slack Variables (for soft constraints):

Variable	Domain	Units	Constraint
$\sigma_{h,k}^{q-}$	≥ 0	m ³ /s	Turbined flow below minimum
$\sigma_{h,k}^{o-}$	≥ 0	m ³ /s	Outflow below minimum
$\sigma_{h,k}^{o+}$	≥ 0	m ³ /s	Outflow above maximum
$\sigma_{h,k}^{g-}$	≥ 0	MW	Generation below minimum
$\sigma_{h,k}^{e+}, \sigma_{h,k}^{e-}$	≥ 0	m ³ /s	Evaporation violation
$\sigma_{h,k}^r$	≥ 0	m ³ /s	Water withdrawal violation
σ_h^{inf}	≥ 0	m ³ /s	Inflow non-negativity (if enabled)

3.4 Dual Variables

Dual variables are essential for cut coefficient computation:

Symbol	Constraint	Cut Coefficient
π_h^{wb}	Water balance	$\beta_h^v = \zeta \cdot \pi_h^{wb}$
$\pi_{h,\ell}^{lag}$	AR lag fixing	$\beta_{h,\ell}^{lag} = \pi_{h,\ell}^{lag}$
$\pi_{b,k}^{lb}$	Load balance	Marginal cost of energy
λ_i	Benders cut i	Cut activity indicator

Sign Convention for Cut Coefficients:

The relationship between LP duals (π) and cut coefficients (β) depends on constraint orientation. For state-linking constraints of the form:

- $x_t = \hat{x}_{t-1} + \dots$ (incoming state on RHS): $\beta = +\pi$
- $x_t - \hat{x}_{t-1} = \dots$ (incoming state on LHS with minus): $\beta = -\pi$

POWE.RS writes the water balance as: $v_h = \hat{v}_h + \zeta[\dots]$ (incoming storage $\hat{v}_h$ on RHS).

With this convention, the cut coefficient is: $\beta_h^v = +\zeta \cdot \pi_h^{wb}$

Intuition: A positive π_h^{wb} means “an extra hm³ of incoming storage reduces cost” (water has value), so the cut should increase with more initial storage, giving a positive β_h^v .

For AR lag constraints: $a_{h,\ell} = \hat{a}_{h,\ell}$ (fixing to incoming value), we have $\beta_{h,\ell}^{lag} = +\pi_{h,\ell}^{lag}$.

4. Base LP Formulation

This section presents the complete stage subproblem LP. The formulation uses **parallel blocks** by default (see Section 5 for chronological blocks variant).

4.0 Cost and Penalty Taxonomy

The objective function includes several cost categories with distinct purposes and typical magnitudes. Understanding this taxonomy is essential for setting appropriate parameter values and interpreting solution reports.

4.0.1 Cost Categories Overview

Category	Purpose	Examples	Typical Magnitude
Resource Costs	Actual generation/operational costs	Thermal fuel, contract prices, pumping energy	\$ 50-500/MWh
Economic Signals	Represent opportunity cost or value	Deficit (load shedding), export revenue	\$ 1,000-10,000/MWh
Regularization Costs	Avoid degenerate solutions, guide solver	Spillage, exchange, excess	\$ 0.001-10/unit
Operational Violation Penalties	Discourage undesirable but feasible operations	Minimum outflow, generation minimum	\$ 500-5,000/unit
Physical Violation Penalties	Discourage physically impossible operations	Negative inflow, storage beyond limits	Very high (\$ 10,000+/unit)

4.0.2 Detailed Cost Definitions

Resource Costs (Actual Operating Expenses) These represent real costs incurred during operation:

Cost	Symbol	Units	Typical Values	Objective Term
Thermal generation	$c_{j,s}^{th}$	\$/MWh	50-500	$\sum_{j,k,s} \tau_k \cdot c_{j,s}^{th} \cdot g_{j,k,s}$
Import contract	c_c^{imp}	\$/MWh	100-300	$\sum_{c,k} \tau_k \cdot c_c^{imp} \cdot \chi_{c,k}^{in}$
Pumping electricity	c_j^{pump}	\$/MWh	Spot price	$\sum_{j,k} \tau_k \cdot c_j^{pump} \cdot \gamma_j \cdot p_{j,k}$

Economic Signals (Opportunity Cost / Value of Lost Load)

Cost	Symbol	Units	Typical Values	Purpose
Deficit (load shedding)	$c_{b,s}^{def}$	\$/MWh	5,000-10,000	Represents value of unserved energy
Export revenue	c_c^{out}	\$/MWh	50-200	Revenue from exports (negative cost)

Regularization Costs (Solution Guidance) These are small costs that prevent degenerate solutions without significantly affecting the optimal policy:

Cost	Symbol	Units	Typical Values	Purpose
Spillage	c_h^{spill}	\$/ (m ³ /s · h)	0.001-0.01	Prefer turbinning over spilling when indifferent
Diversion	c_h^{div}	\$/ (m ³ /s · h)	0.001-0.01	Prefer main channel flow
Exchange	c_ℓ^{exch}	\$/MWh	0.01-1.0	Prevent unnecessary power flows
Excess generation	c_b^{exc}	\$/MWh	0.001-0.1	Eliminate slack generation

Note: Regularization costs should be at least 2-3 orders of magnitude smaller than economic costs to avoid distorting the optimal solution.

Operational Violation Penalties (Soft Constraints)

Penalty	Symbol	Units	Typical Values	Violated Constraint
Turbined flow minimum	c^{q-}	\$/ (m ³ /s · h)	500-1,000	$q_{h,k} \geq \underline{Q}_h$
Outflow minimum	c^{o-}	\$/ (m ³ /s · h)	500-1,000	$o_{h,k} \geq \underline{O}_h$
Outflow maximum	c^{o+}	\$/ (m ³ /s · h)	500-1,000	$o_{h,k} \leq \bar{O}_h$
Generation minimum	c^{g-}	\$/MWh	1,000-2,000	$g_{h,k} \geq \underline{G}_h$

Physical Violation Penalties (Infeasibility Avoidance)

Penalty	Symbol	Units	Typical Values	Purpose
Negative inflow	c^{inf}	\$/ (m ³ /s · h)	10,000+	PAR(p) model produces negative value
Evaporation violation	c^{evap}	\$/ (m ³ /s · h)	5,000+	Computed evaporation exceeds capacity
Withdrawal violation	c^{with}	\$/ (m ³ /s · h)	5,000+	Committed withdrawal cannot be met

4.0.3 Penalty Priority and Hierarchy When setting penalties, ensure the following ordering (from highest to lowest):

1. **Physical violations** (c^{inf} , c^{evap}): Must be prohibitively high to ensure LP feasibility represents physical reality
2. **Deficit**: Represents value of lost load; should exceed any generation cost
3. **Operational violations**: Should exceed typical marginal cost but allow violation when physically necessary
4. **Resource costs**: Market-based or fuel-based
5. **Regularization**: Near-zero to avoid solution distortion

Mathematical Requirement:

$$c^{inf} > c^{def} > c^{q-}, c^{o-}, c^{g-} > c^{th} > c^{spill}, c^{exch}$$

Note on Thermal Plants: Unlike hydro plants, thermal plants do not have slack variables for minimum generation constraints. Thermal bounds ($\underline{G}_j$, $\bar{G}_j$) are treated as hard constraints. If operational requirements (e.g., minimum take-or-pay contracts) cannot be met, the bounds should be adjusted in `thermal_bounds.parquet`. This design choice reflects that thermal dispatch is directly controllable, whereas hydro constraints may be violated due to exogenous inflow uncertainty.

Note on Storage Targets: The `target_storage_hm3` in filling hydros is not a cost term but a constraint target for the dead-volume filling period. See `DATA_MODEL_SPECIFICATION` for filling hydro behavior.

4.0.4 Objective Function Structure The complete stage objective is:

$$\min \underbrace{C^{resource}}_{\text{thermal, contracts}} + \underbrace{C^{deficit}}_{\text{load shedding}} + \underbrace{C^{regularization}}_{\text{spillage, exchange}} + \underbrace{C^{penalty}}_{\text{soft constraints}} + \theta$$

where each component is summed over blocks with appropriate time weighting:

$$C^{component} = \sum_{k \in \mathcal{K}} \tau_k \cdot (\text{cost terms for component})$$

4.1 Objective Function

$$\begin{aligned} \min \sum_{k \in \mathcal{K}} \tau_k \left[\underbrace{\sum_{b \in \mathcal{B}} \sum_{s \in \mathcal{S}_b} c_{b,s}^{def} \delta_{b,k,s}}_{\text{Deficit cost}} + \underbrace{\sum_{b \in \mathcal{B}} c_b^{exc} \epsilon_{b,k}}_{\text{Excess cost}} + \underbrace{\sum_{j \in \mathcal{J}} \sum_s c_{j,s}^{th} g_{j,k,s}}_{\text{Thermal cost}} \right. \\ + \underbrace{\sum_{l \in \mathcal{L}} c_l^{exch} (f_{l,k}^+ + f_{l,k}^-)}_{\text{Exchange cost (regularization)}} + \underbrace{\sum_{h \in \mathcal{H}} c_h^{spill} s_{h,k}}_{\text{Spillage cost (regularization)}} + \underbrace{\sum_{h \in \mathcal{H}} c_h^{div} u_{h,k}}_{\text{Diversion cost (regularization)}} \\ + \underbrace{\sum_{c \in \mathcal{C}^{in}} c_c^{in} c_{c,k}^{in} - \sum_{c \in \mathcal{C}^{out}} c_c^{out} c_{c,k}^{out}}_{\text{Contract cost (import - export revenue)}} + \underbrace{\sum_{j \in \mathcal{P}} c_j^{pump} p_{j,k}}_{\text{Pumping cost}} \\ \left. + \underbrace{\text{Slack penalty terms}}_{\text{See Section 4.0.2}} \right] + \theta \end{aligned}$$

4.2 Load Balance Constraint

For each bus $b \in \mathcal{B}$ and block $k \in \mathcal{K}$:

$$\begin{aligned} \sum_{h \in \mathcal{H}_b} g_{h,k} + \sum_{j \in \mathcal{J}_b} \sum_s g_{j,k,s} + \sum_{l: \text{target}=b} \eta_l f_{l,k}^+ + \sum_{l: \text{source}=b} \eta_l f_{l,k}^- + \sum_{c \in \mathcal{C}_b^{in}} c_{c,k}^{in} \\ - \sum_{l: \text{source}=b} f_{l,k}^+ - \sum_{l: \text{target}=b} f_{l,k}^- - \sum_{c \in \mathcal{C}_b^{out}} c_{c,k}^{out} - \sum_{j \in \mathcal{P}_b} \gamma_j p_{j,k} + \sum_{s \in \mathcal{S}_b} \delta_{b,k,s} - \epsilon_{b,k} = D_{b,k} \end{aligned}$$

Dual variable: $\pi_{b,k}^{lb}$ (marginal cost of energy at bus b , block k)

4.3 Hydro Water Balance

For each hydro $h \in \mathcal{H}$ (parallel blocks formulation):

$$v_h = \hat{v}_h + \zeta \left[a_h + \sum_{k \in \mathcal{K}} w_k \left(\underbrace{\sum_{i \in \mathcal{U}_h} (q_{i,k} + s_{i,k} + u_{i,k})}_{\text{Inflow from upstream}} + \underbrace{\sum_{i: \text{div}=h} u_{i,k}}_{\text{Diverted inflow}} + \underbrace{\sum_{j: \text{dest}=h} p_{j,k}}_{\text{Pumped inflow}} \right. \right. \\ \left. \left. - \underbrace{q_{h,k} - s_{h,k} - u_{h,k}}_{\text{Outflows}} - \underbrace{e_{h,k}}_{\text{Evaporation}} - \underbrace{r_{h,k}}_{\text{Withdrawal}} - \underbrace{\sum_{j: \text{src}=h} p_{j,k}}_{\text{Pumped outflow}} \right) \right]$$

where: - $\hat{v}_h$ = incoming storage (state from previous stage) - a_h = incremental inflow (from AR model, see Section 4.4) - $w_k = \tau_k / \sum_j \tau_j$ = block weight - $\zeta = 0.0036 \times \sum_k \tau_k$ = time conversion factor

Dimensional Consistency: - LHS: v_h [hm^3] - RHS: $\hat{v}_h$ [hm^3] + ζ [$\text{hm}^3/(\text{m}^3/\text{s})$] $\times$ (flow terms [m^3/s]) - The factor ζ converts all flow rates (m^3/s) to volumes (hm^3) accumulated over the stage - Block weights w_k are dimensionless and sum to 1 - The AR inflow a_h is in m^3/s (average rate over the stage)

Dual variable: π_h^{wb} (water value, used for cut coefficients)

4.4 AR Inflow Dynamics

The incremental inflow a_h is determined by the PAR(p) autoregressive model:

$$a_h = \underbrace{\left(\mu_t - \sum_{\ell=1}^{P_h} \psi_\ell \mu_{t-\ell} \right)}_{\text{deterministic base}} + \underbrace{\sum_{\ell=1}^{P_h} \psi_\ell \cdot a_{h,\ell}}_{\text{lag contribution}} + \underbrace{\sigma_t \cdot \eta_t}_{\text{stochastic innovation}}$$

State expansion: To maintain the Markov property, lagged inflows are state variables with fixing constraints:

$$a_{h,\ell} = \hat{a}_{h,\ell} \quad \forall h \in \mathcal{H}, \ell \in \{1, \dots, P_h\}$$

Dual variable: $\pi_{h,\ell}^{lag}$ (value of inflow history, used for cut coefficients)

See Section 8 for the complete PAR(p) model specification.

4.5 Hydro Generation Constraints

Constant Productivity Model (for each hydro $h \in \mathcal{H}^{op}$, block k):

$$g_{h,k} = \rho_h \cdot q_{h,k}$$

FPHA Model (for each plane $m \in \mathcal{M}_h$, hydro h , block k):

$$g_{h,k} \leq \gamma_0^m + \gamma_v^m \cdot v_h^{avg} + \gamma_q^m \cdot q_{h,k} + \gamma_s^m \cdot s_{h,k}$$

where v_h^{avg} is the average storage during the stage (see Section 6 for details).

4.6 Outflow Constraints

Outflow Definition (per hydro h , block k):

$$o_{h,k} = q_{h,k} + s_{h,k}$$

Clarification: Outflow o represents water released to the downstream channel (affecting tailrace level). It does NOT include: - **Withdrawal** $r_{h,k}$: Consumptive use removed from the system (irrigation, water supply) - **Diversion** $u_{h,k}$: Water bypassed to a separate channel (not affecting main tailrace)

The water balance (Section 4.3) accounts for all flows: inflow $-(q + s + u + r)$ - evaporation = storage change.

Outflow Bounds (with slacks for soft enforcement):

$$\underline{Q}_h - \sigma_{h,k}^{o-} \leq o_{h,k} \leq \bar{Q}_h + \sigma_{h,k}^{o+}$$

4.7 Minimum Constraints

Turbined Flow Minimum (per hydro h , block k):

$$q_{h,k} + \sigma_{h,k}^{q-} \geq \underline{Q}_h$$

Generation Minimum (per hydro h , block k):

$$g_{h,k} + \sigma_{h,k}^{g-} \geq \underline{G}_h$$

4.8 Slack Penalties and Soft Constraints

Slack variables allow constraint violations at a cost. The penalty terms in the objective are:

$$\text{Penalties} = \sum_{h,k} \left[c^{q-} \sigma_{h,k}^{q-} + c^{o-} \sigma_{h,k}^{o-} + c^{o+} \sigma_{h,k}^{o+} + c^{g-} \sigma_{h,k}^{g-} + c^e (\sigma_{h,k}^{e+} + \sigma_{h,k}^{e-}) + c^r \sigma_{h,k}^r \right]$$

Penalty Priority (highest to lowest): 1. Stage-specific override (from Parquet files) 2. Entity-specific override (from JSON) 3. Global default (from `penalties.json`)

Typical penalty values:

Constraint	Typical Penalty	Units
Turbined min	500	\$/ (m ³ /s · h)
Outflow min/max	500	\$/ (m ³ /s · h)
Generation min	1000	\$/MWh
Evaporation	5000	\$/ (m ³ /s · h)
Water withdrawal	1000	\$/ (m ³ /s · h)

4.9 Generic Constraints

User-defined linear constraints (per constraint $g \in \mathcal{G}$):

$$\sum_e \gamma_{g,e} \cdot x_e \quad \{\leq, =, \geq\} \quad b_g$$

where x_e can reference any LP variable using expression syntax: - `hydro_storage(id)`, `hydro_turbined(id)`, `hydro_spillage(id)` - `thermal_generation(id)`, `bus_deficit(id)`, etc.

Generic constraints can have optional slack variables with configurable penalties.

4.10 Benders Cuts

For each active cut i from previous iterations:

$$\theta \geq \alpha_i + \sum_{h \in \mathcal{H}} \beta_{i,h}^v \cdot v_h + \sum_{h,\ell} \beta_{i,h,\ell}^{lag} \cdot a_{h,\ell}$$

where: - α_i = cut intercept (RHS) - $\beta_{i,h}^v$ = coefficient for storage state variable - $\beta_{i,h,\ell}^{lag}$ = coefficient for AR lag state variable

Cuts are pre-allocated and toggled active/inactive via bound changes for warm-starting efficiency.

5. Block Formulation Variants

Within each stage, load is divided into **blocks** representing different periods (e.g., peak, off-peak, or hourly resolution). POWE.RS supports two block formulations:

5.1 Parallel Blocks (Default)

In parallel blocks mode, all blocks within a stage are **independent**—there is no intra-stage storage dynamics.

Water Balance (Parallel) A single water balance constraint spans all blocks:

$$v_h = \hat{v}_h + \zeta \left[a_h + \sum_{k \in \mathcal{K}} w_k \cdot \text{net_flows}_{h,k} \right]$$

where: - $w_k = \tau_k / \sum_j \tau_j$ is the block weight - $\text{net_flows}_{h,k}$ = inflows from upstream – outflows – evaporation – withdrawal

This formulation assumes the reservoir can freely redistribute water across blocks within the stage.

Characteristics

Aspect	Description
LP size	Smaller (one water balance per hydro)
Storage dynamics	End-of-stage only
Use case	Long-term strategic planning
Configuration	<code>modeling.block_mode = "parallel"</code>

5.2 Chronological Blocks

In chronological blocks mode, blocks are **sequential** within each stage, enabling modeling of intra-stage storage dynamics (e.g., daily cycling patterns within a monthly stage).

Additional Variables

Variable	Domain	Units	Description
$v_{h,k}$	$[V_h, \bar{V}_h]$	hm ³	Storage at end of block k

The end-of-stage storage (state variable) is: $v_h = v_{h,|\mathcal{K}|}$

Block 1 Water Balance

$$v_{h,1} = \hat{v}_h + \zeta_1 [a_h \cdot w_1 + \text{net_flows}_{h,1}]$$

where $\zeta_1 = 0.0036 \times \tau_1$ is the time conversion for block 1.

Subsequent Blocks Water Balance For $k = 2, \dots, |\mathcal{K}|$:

$$v_{h,k} = v_{h,k-1} + \zeta_k [a_h \cdot w_k + \text{net_flows}_{h,k}]$$

State Variable Definition Only **end-of-stage storage** is a state variable:

$$v_h = v_{h,|\mathcal{K}|}$$

Inter-block storages $v_{h,k}$ for $k < |\mathcal{K}|$ are internal LP variables—not state variables. This ensures: 1. Cuts are computed with respect to end-of-stage storage only 2. State dimension does not increase with number of blocks

Dual Extraction for Cuts For cut generation, we need the dual of block 1’s water balance (containing $\hat{v}_h$):

$$\beta_h^{storage} = \pi_{h,1}^{wb}$$

Characteristics

Aspect	Description
LP size	Larger ($N_{hydro} \times (\mathcal{K} - 1)$ additional variables/constraints)
Storage dynamics	Intra-stage cycling modeled
Use case	Short-term planning with storage cycling
Configuration	<code>modeling.block_mode = "chronological"</code>

5.3 Comparison Summary

Aspect	Parallel Blocks	Chronological Blocks
Water balance	1 per hydro per stage	$ \mathcal{K} $ per hydro per stage
Inter-block storage	Not modeled	Explicit continuity
State variables	End-of-stage only	End-of-stage only
LP variables	Fewer	More
LP constraints	Fewer	More
Intra-stage dynamics	None	Full

6. Hydro Production Function Models

The hydro generation constraint relates turbined flow to electrical output. POWE.RS supports two models:

6.1 Constant Productivity Model

The simplest model assumes linear relationship:

$$g_{h,k} = \rho_h \cdot q_{h,k}$$

where ρ_h (MW per m³/s) is the hydro productivity, typically:

$$\rho_h = \frac{9.81 \times \eta_h \times H_h^{ref}}{1000}$$

with: - η_h = turbine efficiency (typically 0.85-0.92) - H_h^{ref} = reference net head (meters)

Characteristics:

- 1 equality constraint per hydro per block
- Simple, fast
- Ignores head variation with storage

6.2 FPHA (Four-Point Head Approximation)

For accurate modeling of hydroelectric generation, FPHA (Função de Produção Hidrelétrica Aproximada) captures the nonlinear relationship between storage, flow, spillage, and generation through a piecewise-linear approximation.

Notation and Terminology This section uses consistent notation with the LP formulation (Section 4). The following table maps POWE.RS symbols to equivalent CEPEL/Portuguese terminology for practitioners familiar with DECOMP/DESSEM:

POWE.RS	CEPEL/Portuguese	Description	Units
ϕ	FPH	Hydro production function	MW
v	V	Reservoir storage	hm ³
q	Q	Turbined flow	m ³ /s
s	S	Spillage	m ³ /s
g_h	GH	Hydro generation	MW
h_{fore}	h_{mon} (montante)	Forebay (upstream) level	m
h_{tail}	h_{jus} (jusante)	Tailrace (downstream) level	m
h_{net}	h_{liq} (líquida)	Net head	m
h_{loss}	h_{PerdH} (perda hidráulica)	Hydraulic losses	m
q_{out}	Q_{jus}	Total downstream outflow	m ³ /s
q_{lat}	Q_{lat}	Lateral tributary flow	m ³ /s

6.2.1 Exact Production Function The **exact hydroelectric production function** relates generation to the operating state:

$$\phi(v, q, q_{out}) = \rho(q, h_{net}) \times q \times h_{net}$$

where: - v = reservoir storage volume (hm³) - q = turbined flow (m³/s) - q_{out} = total downstream outflow affecting tailrace level (m³/s) - h_{net} = net head (m) - ρ = specific productivity (MW · s/m⁴)

The **net head** is computed as:

$$h_{net}(v, q, q_{out}) = h_{fore}(v) - h_{tail}(q_{out}) - h_{loss}(q)$$

where: - $h_{fore}(v)$ = forebay (upstream reservoir) level as function of storage - $h_{tail}(q_{out})$ = tailrace (downstream channel) level as function of total outflow - $h_{loss}(q)$ = hydraulic head losses in penstock and turbines

Why linearization is needed: The production function ϕ is nonlinear in (v, q) due to: 1. The product $q \times h_{net}$ (bilinear term) 2. Nonlinear topology functions $h_{fore}(v)$ and $h_{tail}(q_{out})$ 3. Flow-dependent hydraulic losses

For LP formulation, we approximate ϕ with a set of linear hyperplanes.

6.2.2 Topology Functions (POWE.RS Approach) Unlike CEPEL models which use 4th-degree polynomial fits, POWE.RS uses **tabular data with linear interpolation** for topology functions. This approach is: - More transparent (no hidden polynomial artifacts) - Easier to validate against surveyed data - Flexible for any reservoir geometry

Forebay Level $h_{fore}(v)$ The upstream water level is obtained from `hydro_geometry.parquet`:

volume_hm3	height_m	area_km2
v_1	h_1	A_1
v_2	h_2	A_2
...	...	...

Interpolation: For storage v where $v_i \leq v < v_{i+1}$:

$$h_{fore}(v) = h_i + \frac{h_{i+1} - h_i}{v_{i+1} - v_i} \times (v - v_i)$$

Tailrace Level $h_{tail}(q_{out})$ The downstream water level depends on total outflow. From `hydro_production_data.parquet`:

Polynomial model (CEPEL-compatible):

$$h_{tail}(q_{out}) = c_0 + c_1 q_{out} + c_2 q_{out}^2 + c_3 q_{out}^3 + c_4 q_{out}^4$$

Piecewise-linear model (POWE.RS native): | outflow_m3s | tailrace_m | |-----|-----| | q_1 | $h_{tail,1}$ | | q_2 | $h_{tail,2}$ | | ... | ... |

With linear interpolation between points.

Total downstream flow: $q_{out} = q + s + q_{lat}$ where: - q = turbined flow - s = spillage - q_{lat} = lateral inflows between reservoir and tailrace (from tributaries or upstream plants)

Hydraulic Losses $h_{loss}(q)$ Two models are supported:

Factor model (proportional to gross head):

$$h_{loss}(q) = k_{loss} \times (h_{fore} - h_{tail})$$

where k_{loss} is typically 0.01-0.05 (1-5% losses).

Constant model (fixed head loss):

$$h_{loss}(q) = \Delta h_{const}$$

where Δh_{const} is in meters (typically 1-5m).

Flow-dependent model (future extension):

$$h_{loss}(q) = k_q \times q^2$$

This captures the quadratic friction losses in penstocks.

6.2.3 Variable Productivity Model The **specific productivity** converts hydraulic power to electrical power:

$$\rho(q, h_{net}) = \frac{g \times \eta(q)}{1000}$$

where: - $g = 9.81 \text{ m/s}^2$ (gravitational acceleration) - $\eta(q)$ = turbine-generator efficiency (dimensionless, typically 0.85-0.93) - Factor 1000 converts W to kW

The full generation formula becomes:

$$g_h = \frac{9.81 \times \eta \times q \times h_{net}}{1000} \quad [\text{MW}]$$

Constant efficiency (current implementation):

$$\eta(q) = \eta_{ref} \quad (\text{constant})$$

Variable efficiency (future extension, see Appendix C):

$$\eta(q) = \eta_{max} \times f\left(\frac{q}{q_{nom}}\right)$$

where f is a characteristic curve peaking near nominal flow.

Reference Productivity For the constant productivity model, the reference value is:

$$\rho_{ref} = \frac{9.81 \times \eta_{ref} \times h_{ref}}{1000} \quad [\text{MW per m}^3/\text{s}]$$

where h_{ref} is the reference net head (typically at 65% storage).

6.2.4 Hyperplane Fitting Algorithm POWERS supports two approaches for obtaining FPHA hyperplanes:

1. **Pre-fitted:** Read coefficients from `fpha_hyperplanes.parquet` 2. **Computed:** Generate from topology data during preprocessing

Algorithm: FPHA Hyperplane Fitting **Input:**

- Topology data: $h_{fore}(v)$, $h_{tail}(q_{out})$, $h_{loss}(q)$
- Operating bounds: $[v_{min}, v_{max}]$, $[0, q_{max}]$
- Discretization: n_v volume points, n_q turbine flow points
- Reference spillage: s_{ref} (typically 0 or average expected spillage)

Output:

- Set of hyperplanes $\{(\gamma_0^m, \gamma_v^m, \gamma_q^m, \gamma_s^m)\}_{m=1}^M$
- Correction factor κ (note: we use κ to avoid collision with cut intercept α)

Algorithm: FPHA_Fit

1. **DISCRETIZE operating window**

- Create volume grid: $v_{grid} = \text{linspace}(v_{min}, v_{max}, n_v)$
- Create flow grid: $q_{grid} = \text{linspace}(0, q_{max}, n_q)$
- Stage-dependent configuration:
 - Near-term stages: higher resolution ($n_v = 7$, $n_q = 15$)
 - Far-future stages: lower resolution ($n_v = 3$, $n_q = 5$)

2. **EVALUATE exact production function at grid points**

For each (v_i, q_j) in grid, compute:

- Forebay head: $h_{fore} = \text{interpolate}(\text{geometry_table}, v_i)$

- Total outflow: $q_{out} = q_j + s_{ref}$
- Tailrace head: $h_{tail} = \text{interpolate}(\text{tailrace_table}, q_{out})$
- Head loss: $h_{loss} = \text{compute_loss}(q_j, h_{fore}, h_{tail})$
- Net head: $h_{net} = h_{fore} - h_{tail} - h_{loss}$
- Generation: If $h_{net} > 0$, then $g_{exact}[i, j] = \rho \times q_j \times h_{net}$, else $g_{exact}[i, j] = 0$

3. BUILD convex hull of generation surface

- Create 3D point cloud:
 - points = $\{(v_i, q_j, g_{exact}[i, j]) \mid \forall (i, j) \text{ with } g > 0\}$
- Compute upper convex hull (concave envelope where generation $\leq$ surface):
 - Run qhull: `hull = qhull(points, options = "Qt Qc")`
- Extract facets with downward-facing normals (upper hull):
 - Initialize planes = []
 - For each facet in `hull.facets`:
 - * If `facet.normal[2] < 0` (upward in g_h direction):
 - Extract plane coefficients $(\gamma_0, \gamma_v, \gamma_q)$ from facet
 - Append to planes

4. COMPUTE correction factor κ

Apply correction to ensure $\text{FPHA} \leq \phi$ everywhere:

- Initialize $\kappa = 1.0$
- For each (v_i, q_j) in grid:
 - Compute $g_{fpha} = \max_m \{\gamma_0^m + \gamma_v^m \cdot v_i + \gamma_q^m \cdot q_j\}$
 - If $g_{exact}[i, j] > 0$ AND $g_{fpha} > 0$:
 - * Update $\kappa = \min(\kappa, g_{exact}[i, j] / g_{fpha})$
- Scale all intercepts: $\gamma_0^m = \kappa \times \gamma_0^m$ for each plane m

Optional MSE minimization: $\kappa = \arg \min_{\kappa} \sum_{i,j} (\kappa \cdot g_{fpha}[i, j] - g_{exact}[i, j])^2$

5. ADD spillage dimension (secant approximation)

Spillage affects tailrace level, reducing net head:

- Compute tailrace sensitivity:

$$\frac{dh_{tail}}{ds} = \frac{h_{tail}(q_{ref} + s_{ref} + \Delta s) - h_{tail}(q_{ref} + s_{ref})}{\Delta s}$$

- For each plane m , add spillage coefficient:

$$\gamma_s^m = -\rho \times q_{ref} \times \frac{dh_{tail}}{ds}$$

6. RETURN planes and metadata

- **planes:** $\{(\gamma_0^m, \gamma_v^m, \gamma_q^m, \gamma_s^m) \mid m = 1, \dots, M\}$
- **kappa:** κ
- **num_planes:** M
- **fitting_bounds:** $\{v_{min}, v_{max}, q_{max}\}$
- **grid_resolution:** $\{n_v, n_q\}$

Qhull Algorithm Details The qhull library (or equivalent) computes convex hulls in $\mathbb{R}^n$. For FPHA fitting:

1. **Input points:** $(v_i, q_j, g_{i,j})$ for all grid points
2. **Convex hull:** Find the minimal convex polytope containing all points
3. **Upper hull extraction:** Keep only facets where generation is maximized (the “roof” of the polytope)
4. **Plane coefficients:** Each facet defines a half-space $g_h \leq \gamma_0 + \gamma_v v + \gamma_q q$

Implementation options:

- Rust: Use `convex_hull` from computational geometry crate
- External: Call qhull via FFI or subprocess
- Simplified: Use Delaunay triangulation and filter upper facets

6.2.5 Correction Factor Calculation The correction factor κ ensures the approximation is conservative (never overestimates generation).

Notation Note: We use κ (kappa) for the FPHA correction factor to avoid collision with α , which is used for Benders cut intercepts throughout this document (see Section 10).

Worst-Case Approach (Default)

$$\kappa = \min_{(v,q) \in \text{grid}} \left\{ \frac{\phi(v,q)}{\max_m (\gamma_0^m + \gamma_v^m v + \gamma_q^m q)} \right\}$$

This guarantees $g_{h,FPHA} \leq \phi$ everywhere in the operating region.

MSE Minimization Approach

$$\kappa = \arg \min_{\kappa} \sum_{(v_i, q_j)} (\kappa \cdot g_{FPHA}(v_i, q_j) - \phi(v_i, q_j))^2$$

Closed-form solution:

$$\kappa = \frac{\sum_{i,j} g_{FPHA} \cdot \phi}{\sum_{i,j} g_{FPHA}^2}$$

Typical Values

Reservoir Type	Typical κ	Notes
High-head storage	0.97-0.99	Significant head variation
Medium-head	0.98-1.00	Moderate approximation error
Run-of-river	0.99-1.00	Nearly constant head

6.2.6 Spillage and Lateral Flow Effects

Downstream Level Dependency The tailrace level depends on total flow through the downstream channel:

$$q_{out} = q + s + q_{lat}$$

where q_{lat} includes: - Lateral tributaries entering between dam and tailrace - Outflow from upstream plants in cascade - Return flows from irrigation or other withdrawals

Secant Approximation for Spillage Since spillage s affects tailrace level, it indirectly affects generation. The FPHA constraint incorporates this through γ_s :

$$\gamma_s^m = -\rho \times q_{ref} \times \left. \frac{\partial h_{tail}}{\partial q_{out}} \right|_{q_{out,ref}}$$

Physical interpretation: Each additional m^3/s of spillage raises the tailrace by $\partial h_{tail}/\partial q_{out}$ meters, reducing net head and thus generation.

Sign convention: $\gamma_s^m < 0$ because spillage reduces generation capacity.

Cascade Effects (Advanced) In cascade systems, upstream spillage affects downstream tailrace levels with a time delay. This creates cross-plant coupling not captured in standard FPHA. For now, POWERS assumes: - Each plant's FPHA is independent - Cascade effects are captured through average expected flows - Future enhancement could add cross-plant correction terms

6.2.7 LP Integration

Final FPHA Constraint For each hydro h , block k , and plane $m \in \mathcal{M}_h$:

$$g_{h,k}^{hy} \leq \kappa \times (\gamma_0^m + \gamma_v^m \cdot v_h^{avg} + \gamma_q^m \cdot q_{h,k} + \gamma_s^m \cdot s_{h,k})$$

Or equivalently with pre-scaled coefficients:

$$g_{h,k}^{hy} \leq \tilde{\gamma}_0^m + \gamma_v^m \cdot v_h^{avg} + \gamma_q^m \cdot q_{h,k} + \gamma_s^m \cdot s_{h,k}$$

where $\tilde{\gamma}_0^m = \kappa \times \gamma_0^m$.

Average Storage Computation The average storage v_h^{avg} over the stage depends on configuration:

Option A: Simple Average (Default)

$$v_h^{avg} = \frac{\hat{v}_h + v_h}{2}$$

where $\hat{v}_h$ is incoming storage and v_h is end-of-stage storage.

Option B: Block-Weighted Average

$$v_h^{avg} = \sum_k w_k \cdot v_{h,k}^{mid}$$

where w_k is the block weight (duration fraction) and $v_{h,k}^{mid}$ is the mid-block storage.

Generation as Independent Variable When using FPHA, the generation variable $g_{h,k}^{hy}$ is **not** directly computed from turbined flow. Instead:

1. Generation is a free LP variable bounded by $[0, \bar{G}_h]$
2. FPHA constraints (one per plane) provide upper bounds
3. The optimizer maximizes generation subject to FPHA constraints
4. At optimum, generation “touches” one of the FPHA planes

Key insight: Because minimizing cost includes maximizing hydro generation (which has zero fuel cost), the optimizer naturally pushes generation to the FPHA surface boundary.

Slack Variables for Soft Constraints For numerical robustness, FPHA constraints can include slack variables:

$$g_{h,k}^{hy} - \sigma_{h,k,m}^{fpha} \leq \tilde{\gamma}_0^m + \gamma_v^m \cdot v_h^{avg} + \gamma_q^m \cdot q_{h,k} + \gamma_s^m \cdot s_{h,k}$$

where $\sigma_{h,k,m}^{fpha} \geq 0$ with high penalty cost. This allows the LP to remain feasible even if operating outside the FPHA validity region.

6.2.8 Water Value and Benders Cuts The FPHA formulation affects how water values are computed and propagated through Benders cuts.

Dual Variables Let π_m^{fpha} be the dual variable for FPHA constraint m . At optimum:

$$\frac{\partial \mathcal{L}}{\partial g_{h,k}^{hy}} = -c_k^{deficit} + \sum_m \pi_m^{fpha} = 0$$

where $c_k^{deficit}$ is the marginal cost of deficit in block k .

Water Value Derivation The marginal value of storage v_h includes the FPHA contribution:

$$\frac{\partial \text{Cost}}{\partial v_h} = \underbrace{\pi_h^{balance}}_{\text{direct value}} + \underbrace{\sum_m \pi_m^{fpha} \cdot \gamma_v^m}_{\text{FPHA contribution}}$$

For the hydro balance constraint with dual $\pi_h^{balance}$:

$$v_h = \hat{v}_h + a_h - q_h - s_h - w_h - e_h$$

Cut Coefficient for Storage The Benders cut coefficient for storage state variable $\hat{v}_h$ is:

$$\beta_{\hat{v}_h} = \pi_h^{balance} + \frac{1}{2} \sum_m \pi_m^{fpha} \cdot \gamma_v^m$$

The factor $\frac{1}{2}$ appears because $v^{avg} = (\hat{v}_h + v_h)/2$, so $\partial v^{avg} / \partial \hat{v}_h = 1/2$.

Model Transition Considerations When a hydro transitions between production models across stages:

Transition	Cut Interpretation	Action
Constant $\rightarrow$ FPHA	Cuts at stage t use constant model	Cut valid but conservative
FPHA $\rightarrow$ Constant	Stage $t + 1$ backward pass uses constant	May overestimate value
FPHA $\rightarrow$ FPHA (different params)	Parameters change	Cuts remain valid if conservative

Recommendation: When using stage-dependent FPHA configuration, ensure the FPHA at stage t is at least as conservative as stage $t + 1$ for cut validity.

6.2.9 Stage-Dependent Configuration POWERS allows different FPHA configurations for different stages, enabling a trade-off between accuracy and computational efficiency.

Fitting Window Selection The operating range $[v_{min}, v_{max}]$ for FPHA fitting can be stage-dependent:

Near-term stages (0-24):

- Use full storage range: $[v_{min}^{phys}, v_{max}^{phys}]$
- Higher resolution: $n_v = 7, n_q = 15$
- All operating scenarios possible

Medium-term stages (25-60):

- Narrower range based on expected operation: $[v_{10\%}, v_{90\%}]$
- Medium resolution: $n_v = 5, n_q = 10$
- Focus on likely operating region

Far-future stages (61+):

- Conservative range centered on equilibrium: $[v_{25\%}, v_{75\%}]$
- Lower resolution: $n_v = 3, n_q = 5$
- Prioritize computational efficiency

Configuration Schema

```
{
  "hydro_id": 42,
  "stage_ranges": [
    {
      "start_stage_id": 0,
      "end_stage_id": 24,
      "model": "fpha",
      "fpha_config": {
        "source": "computed",
        "volume_discretization_points": 7,
        "turbine_discretization_points": 15,
        "fitting_window": {
          "volume_min_hm3": null,
          "volume_max_hm3": null,
          "volume_min_percentile": null,
          "volume_max_percentile": null
        }
      }
    },
    {
      "start_stage_id": 25,
      "end_stage_id": 60,
      "model": "fpha",
      "fpha_config": {
        "source": "computed",
        "volume_discretization_points": 5,
        "turbine_discretization_points": 10,
        "fitting_window": {
          "volume_min_percentile": 10,
          "volume_max_percentile": 90
        }
      }
    },
    {
      "start_stage_id": 61,
      "end_stage_id": null,
      "model": "constant_productivity"
    }
  ]
}
```

Training vs Simulation Phases Future enhancement: Different FPHA configurations for training (SDDP iterations) and simulation (policy evaluation):

- **Training:** Wider fitting windows to ensure cut validity across all scenarios
- **Simulation:** Tighter fitting windows based on observed training trajectories

Note: This requires careful handling to ensure cuts remain valid. See Appendix C for deferred features.

6.3 Linearized Head Model

An intermediate model between constant productivity and full FPHA:

$$g_{h,k}^{hy} = \rho_{ref} \cdot q_{h,k} \cdot (k_0 + k_V \cdot v_h^{avg})$$

where: - k_0, k_V are linearization coefficients derived from $h_{mon}(V)$ - $k_0 = 1 - k_V \cdot V_{ref}$ (normalization at reference

$$\text{volume}) - k_V = \frac{1}{H_{ref}} \cdot \frac{dh_{mon}}{dV} \Big|_{V_{ref}}$$

Characteristics:

- Single constraint (bilinear approximation)
- Captures first-order head variation with storage
- Does not capture spillage effects
- Suitable for medium-term stages

6.4 Model Selection Guidelines

Scenario	Recommended Model	Rationale
High-head storage reservoirs	FPHA	Significant head variation (>20%)
Large storage variation plants	FPHA	Operating across wide volume range
Run-of-river plants	Constant productivity	Nearly constant head
Initial algorithm testing	Constant productivity	Fast iteration, debug focus
Production studies (near-term)	FPHA	Accuracy for operational decisions
Production studies (far-future)	Constant or linearized	Computational efficiency
Post-optimization validation	Compare all models	Verify approximation quality

6.5 FPHA Data Requirements Summary

Data Source	Required Fields	Used For
hydro_geometry.parquet	volume_hm3, height_m	$h_{mon}(V)$ interpolation
hydro_production_data.parquet	tailrace_coeffs or table	$h_{jus}(Q_{jus})$ computation
hydro_production_data.parquet	hydraulic_loss_type, value	$h_{PerdH}(Q)$ computation
hydros.json	productivity_mw_per_m3s	Reference ρ_{ref}
fpha_hyperplanes.parquet	gamma_0, gamma_v, gamma_q, gamma_s	Pre-fitted planes (optional)
hydro_production_models.json	fpha_config per stage	Fitting configuration

7. Equipment-Specific Formulations

This section details the LP constraints for each equipment type.

7.1 Thermal Plants

7.1.1 Standard Thermals Thermal generation uses piecewise-linear cost functions with segments:

Decision Variables:

- $g_{j,k,s}$ = generation at thermal j , block k , cost segment s

Constraints:

Total generation:

$$g_{j,k} = \sum_s g_{j,k,s}$$

Segment bounds:

$$0 \leq g_{j,k,s} \leq \bar{g}_{j,s} \quad \forall s$$

Objective Contribution:

$$\sum_k \tau_k \sum_s c_{j,s}^{th} \cdot g_{j,k,s}$$

Note: POWE.RS does not include binary commitment variables. The model uses continuous relaxation with min/max bounds. For detailed unit commitment, post-process SDDP results with a commitment model.

7.1.2 GNL Thermals (DEFERRED) GNL (Liquefied Natural Gas) plants require dispatch anticipation due to fuel ordering lead times. See [Appendix C](#) for planned formulation.

7.2 Transmission Lines

Decision Variables:

- $f_{l,k}^+$ = direct flow (source $\rightarrow$ target)
- $f_{l,k}^-$ = reverse flow (target $\rightarrow$ source)

Bounds:

$$0 \leq f_{l,k}^+ \leq \bar{F}_l^+, \quad 0 \leq f_{l,k}^- \leq \bar{F}_l^-$$

Load Balance Contribution:

At source bus:

$$-f_{l,k}^+ + \eta_l f_{l,k}^-$$

At target bus:

$$\eta_l f_{l,k}^+ - f_{l,k}^-$$

where $\eta_l = 1 - \text{losses_percent}/100$ accounts for transmission losses.

Objective Contribution:

$$\sum_k \tau_k \cdot c_l^{exch} (f_{l,k}^+ + f_{l,k}^-)$$

Note on Exchange Cost: The cost c_l^{exch} is a **regularization term**, not an actual transmission cost. Its purpose is to: 1. **Prevent degenerate solutions:** Without this term, multiple equivalent solutions exist with different flow patterns 2. **Guide the solver:** Small positive cost encourages minimal power transfers when indifferent 3. **Improve numerical stability:** Reduces cycling in LP simplex iterations

Typical values are very small (\$0.01-1.00/MWh), several orders of magnitude below generation costs. If this cost significantly affects dispatch decisions, the value is set too high.

See Section 4.0.2 (Regularization Costs) for the full taxonomy of penalty vs. cost types.

7.3 Import/Export Contracts

Decision Variables:

- $\chi_{c,k}^{in}$ = import power from contract c
- $\chi_{c,k}^{out}$ = export power to contract c

Bounds:

$$0 \leq \chi_{c,k}^{in} \leq \bar{C}_c, \quad 0 \leq \chi_{c,k}^{out} \leq \bar{C}_c$$

Load Balance Contribution:

At connected bus: $+\chi_{c,k}^{in} - \chi_{c,k}^{out}$

Objective Contribution:

$$\sum_k \tau_k \left(c_c^{imp} \cdot \chi_{c,k}^{in} - c_c^{exp} \cdot \chi_{c,k}^{out} \right)$$

Note: Export revenue is typically positive, hence subtracted from cost.

7.4 Pumping Stations

Pumping stations transfer water from source hydro to destination hydro, consuming electrical power.

Decision Variables:

- $p_{j,k}$ = pumped water flow (m³/s)

Power Consumption:

$$P_{j,k}^{pump} = \gamma_j \cdot p_{j,k}$$

where γ_j is the power consumption rate (MW per m³/s).

Water Balance Impact:

- Source hydro: $-p_{j,k}$ (water removed)
- Destination hydro: $+p_{j,k}$ (water added)

Load Balance Impact: At connected bus: $-P_{j,k}^{pump}$ (power consumed)

Objective Contribution:

$$\sum_k \tau_k \cdot c_j^{pump} \cdot p_{j,k}$$

7.5 Batteries (DEFERRED)

Battery energy storage systems with charge/discharge dynamics. See [Appendix C](#) for planned formulation.

7.6 Non-Controllable Sources (DEFERRED)

Wind and solar generation with stochastic availability. See [Appendix C](#) for planned formulation.

Part III: Stochastic Modeling

8. PAR(p) Inflow Model

8.1 PAR(p) Model Definition

The **Periodic Autoregressive model of order p** (PAR(p)) captures temporal correlation in inflow time series while accounting for seasonal variation in parameters. For hydro h at stage t corresponding to season $m(t)$:

$$a_{h,t} = \mu_{m(t)} + \sum_{\ell=1}^p \psi_{m(t),\ell} (a_{h,t-\ell} - \mu_{m(t-\ell)}) + \sigma_{m(t)} \cdot \varepsilon_t$$

where: - $a_{h,t}$: Incremental inflow at stage t (m³/s) - $\mu_{m(t)}$: Seasonal mean for season $m(t)$ - $\psi_{m(t),\ell}$: Autoregressive coefficient for lag ℓ in season $m(t)$ - $\sigma_{m(t)}$: Seasonal standard deviation of residuals - $\varepsilon_t \sim \mathcal{N}(0,1)$: Innovation (standardized noise) - $m(t)$: Season/period index for stage t (e.g., month 1-12)

8.2 Notation for Fitting

Let $Y_m = \{a_{h,t} : m(t) = m\}$ be the historical observations for season m . Define:

Symbol	Description
N_m	Number of observations for season m

Symbol	Description
$\bar{a}_m$	Sample mean for season m
s_m	Sample standard deviation for season m
$\gamma_m(\ell)$	Autocovariance at lag ℓ for season m
$\rho_m(\ell)$	Autocorrelation at lag ℓ for season m

8.3 Step 1: Seasonal Means and Standard Deviations

Seasonal Mean:

$$\hat{\mu}_m = \bar{a}_m = \frac{1}{N_m} \sum_{t:m(t)=m} a_{h,t}$$

Seasonal Standard Deviation:

$$\hat{s}_m = \sqrt{\frac{1}{N_m - 1} \sum_{t:m(t)=m} (a_{h,t} - \bar{a}_m)^2}$$

8.4 Step 2: Seasonal Autocorrelations

The autocorrelation at lag ℓ for season m is computed from standardized deviations:

Cross-seasonal autocovariance:

For observations at season m with lag ℓ reaching back to season $m - \ell \pmod{M}$, where M is the cycle length):

$$\hat{\gamma}_m(\ell) = \frac{1}{N_m - 1} \sum_{t:m(t)=m} (a_{h,t} - \bar{a}_m) (a_{h,t-\ell} - \bar{a}_{m-\ell})$$

Autocorrelation:

$$\hat{\rho}_m(\ell) = \frac{\hat{\gamma}_m(\ell)}{\hat{s}_m \cdot \hat{s}_{m-\ell}}$$

where $\hat{s}_{m-\ell}$ is the standard deviation of season $m - \ell$ (cyclically, so season 0 = season M).

8.5 Step 3: Yule-Walker Equations

For each season m , the PAR(p) coefficients $\psi_{m,1}, \dots, \psi_{m,p}$ are found by solving the **Yule-Walker system**:

$$\begin{pmatrix} 1 & \hat{\rho}_{m-1}(1) & \hat{\rho}_{m-2}(2) & \cdots & \hat{\rho}_{m-p+1}(p-1) \\ \hat{\rho}_m(1) & 1 & \hat{\rho}_{m-1}(1) & \cdots & \hat{\rho}_{m-p+2}(p-2) \\ \hat{\rho}_m(2) & \hat{\rho}_{m-1}(1) & 1 & \cdots & \hat{\rho}_{m-p+3}(p-3) \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \hat{\rho}_m(p-1) & \hat{\rho}_{m-1}(p-2) & \hat{\rho}_{m-2}(p-3) & \cdots & 1 \end{pmatrix} \begin{pmatrix} \psi_{m,1}^* \\ \psi_{m,2}^* \\ \psi_{m,3}^* \\ \vdots \\ \psi_{m,p}^* \end{pmatrix} = \begin{pmatrix} \hat{\rho}_m(1) \\ \hat{\rho}_m(2) \\ \hat{\rho}_m(3) \\ \vdots \\ \hat{\rho}_m(p) \end{pmatrix}$$

In matrix notation: $\mathbf{R}_m \psi_m^* = r_m$

where: - $\mathbf{R}_m$ is the $p \times p$ correlation matrix (Toeplitz-like but with cross-seasonal correlations) - $r_m = (\hat{\rho}_m(1), \dots, \hat{\rho}_m(p))^T$ is the vector of target autocorrelations

Solution:

$$\hat{\psi}_m^* = \mathbf{R}_m^{-1} r_m$$

8.6 Step 4: Convert to Original Units

The Yule-Walker solution $\psi_{m,\ell}^*$ is for standardized variables. Convert back to original units:

$$\hat{\psi}_{m,\ell} = \psi_{m,\ell}^* \cdot \frac{\hat{s}_m}{\hat{s}_{m-\ell}}$$

8.7 Step 5: Residual Standard Deviation

The residual variance for season m is:

$$\hat{\sigma}_m^2 = \hat{s}_m^2 \left(1 - \sum_{\ell=1}^p \psi_{m,\ell}^* \cdot \hat{\rho}_m(\ell) \right)$$

The residual standard deviation:

$$\hat{\sigma}_m = \hat{s}_m \sqrt{1 - r_m^\top \mathbf{R}_m^{-1} r_m}$$

8.8 Complete PAR(p) Parameter Set

For each hydro h and each season $m \in \{1, \dots, M\}$ (e.g., $M = 12$ for monthly, $M = 52$ for weekly):

Parameter	Formula	Description
μ_m	$\bar{a}_m$	Seasonal mean
$\psi_{m,1}, \dots, \psi_{m,p}$	Yule-Walker solution	AR coefficients
σ_m	$\hat{\sigma}_m$	Residual standard deviation

8.9 Model Order Selection

The PAR order p can vary by season. Common selection criteria:

1. AIC (Akaike Information Criterion):

$$\text{AIC}_m(p) = N_m \ln(\hat{\sigma}_m^2) + 2p$$

2. BIC (Bayesian Information Criterion):

$$\text{BIC}_m(p) = N_m \ln(\hat{\sigma}_m^2) + p \ln(N_m)$$

3. Coefficient significance: Include lag ℓ only if $|\hat{\psi}_{m,\ell}| > 2/\sqrt{N_m}$

8.10 CEPEL PAR(p)-A Variant (Future Extension)

CEPEL's PAR(p)-A model (referenced in Rel-1941_2021) uses: - **Order constraint:** Maximum AR order often fixed at 12 (annual cycle) - **Stationarity enforcement:** Coefficients adjusted to ensure $\sum_{\ell} \psi_{m,\ell} < 1$ - **Lognormal transformation:** Working with $\ln(a_{h,t})$ for strictly positive inflows - **Regional correlation:** Cross-correlation between hydros in same river basin

This variant is not currently implemented but the data model supports it via the `ar_order` and `ar_coef_*` columns in `inflow_models.parquet`.

8.11 Validation Checks

After fitting, verify:

1. **Positive residual variance:** $\hat{\sigma}_m^2 > 0$ for all seasons
2. **Stationarity:** Roots of $1 - \sum_{\ell} \psi_{m,\ell} z^\ell = 0$ lie outside unit circle
3. **Correlation matrix positive definite:** $\mathbf{R}_m$ is invertible
4. **No systematic bias:** Residuals ε_t have mean near zero

9. Inflow Non-Negativity Solution Methods

9.1 Problem Statement

The PAR(p) model can generate negative inflow realizations:

$$a_h = \underbrace{\mu_m - \sum_{\ell=1}^p \psi_{\ell} \mu_{m-\ell}}_{\text{deterministic base}} + \underbrace{\sum_{\ell=1}^p \psi_{\ell} \cdot \hat{a}_{h,\ell}}_{\text{lag contribution}} + \underbrace{\sigma_m \cdot \eta}_{\text{noise term}}$$

When η is sufficiently negative (e.g., $\eta < -2$), the total can become negative, which is physically impossible.

9.2 Method 1: None (`sem_relaxacao`)

Description: No treatment. Negative inflows pass directly to the LP.

LP Formulation: Standard AR constraint (unchanged):

$$a_h = \text{deterministic_base} + \sum_{\ell=1}^p \psi_{\ell} \cdot a_{h,\ell} + \sigma_m \cdot \eta$$

Implications: - LP may become **infeasible** when $a_h < 0$ causes water balance violation - Useful only for debugging or when AR model guarantees positive outputs - **Not recommended for production**

9.3 Method 2: Penalty (`penalizacao`)

Description: Add a slack variable to ensure LP feasibility, with penalty in objective.

Additional Variables:

Variable	Domain	Units	Description
σ_h^{inf}	≥ 0	m ³ /s	Inflow non-negativity slack

Modified AR Constraint:

$$a_h + \sigma_h^{inf} = \text{deterministic_base} + \sum_{\ell=1}^p \psi_{\ell} \cdot a_{h,\ell} + \sigma_m \cdot \eta$$

Interpretation: When the AR model produces negative a_h , the slack σ_h^{inf} absorbs the violation, making the effective inflow:

$$a_h^{effective} = a_h + \sigma_h^{inf} \geq 0$$

Objective Function Addition:

$$+ \sum_{h \in \mathcal{H}} c^{inf} \cdot \sigma_h^{inf} \cdot \zeta$$

where c^{inf} is the penalty cost (default: 1000 \$/(m³/s · h)) and ζ is the time conversion factor.

Advantages: - LP always feasible - Clear cost signal for negative inflow events - Preserves AR dynamics for positive realizations

Disadvantages: - Adds variables and constraints - Slightly affects marginal water values

Recommended for most production cases.

9.4 Method 3: Truncation (**truncamento**)

Description: Hard truncation of negative values to zero during scenario generation.

Scenario Generation:

$$a_h = \max \left(0, \text{deterministic_base} + \sum_{\ell=1}^p \psi_{\ell} \cdot \hat{a}_{h,\ell} + \sigma_m \cdot \eta \right)$$

LP Formulation: Standard AR constraint with the already-truncated a_h value:

$$a_h = (\text{truncated value from scenario})$$

Advantages: - Simple implementation - No additional LP variables - Fast computation

Disadvantages: - **Biases the distribution:** Shifts mean upward - **Breaks AR dynamics:** When truncation occurs, temporal correlation is disrupted - May affect long-term storage dynamics

9.5 Method 4: Truncation with Penalty (**truncamento_penalizacao**)

Description: Hybrid approach that truncates the final inflow but penalizes the statistical violation in the noise term. Based on the YP_FINF slack in SPARHTACUS/SPTcpp.

Additional Variables:

Variable	Domain	Units	Description
ξ_h	≥ 0	-	Noise adjustment slack (dimensionless)

Modified AR Constraint (in two parts):

Part A - Modified noise term:

$$\eta_h^{adj} = \eta_h + \xi_h$$

where η_h is the original (possibly very negative) noise realization.

Part B - Inflow with adjusted noise:

$$a_h = \text{deterministic_base} + \sum_{\ell=1}^p \psi_{\ell} \cdot a_{h,\ell} + \sigma_m \cdot \eta_h^{adj}$$

Non-negativity constraint:

$$a_h \geq 0$$

Interpretation: The optimizer chooses ξ_h to be the minimum adjustment needed to make $a_h \geq 0$:

$$\xi_h = \max \left(0, -\eta_h - \frac{\text{deterministic_base} + \sum_{\ell} \psi_{\ell} \cdot \hat{a}_{h,\ell}}{\sigma_m} \right)$$

Objective Function Addition:

$$+ \sum_{h \in \mathcal{H}} c^{inf} \cdot \sigma_m \cdot \xi_h \cdot \zeta$$

The penalty is proportional to $\sigma_m \cdot \xi_h$, which is the actual inflow adjustment in m³/s.

Advantages: - Preserves AR model structure better than pure truncation - Penalty signals statistical violation severity - Effective inflow is always non-negative

Disadvantages: - More complex formulation - Requires careful interaction with noise generation

9.6 Comparison Summary

Method	LP Size	Bias	AR Preservation	Feasibility	Recommendation
none	Base	None	Full	May fail	Debugging only
penalty	+vars/cons	Minimal	Full	Guaranteed	Production
truncation	Base	Upward	Partial	Guaranteed	Quick studies
truncation_with_penalty	+vars/cons	Minimal	Full	Guaranteed	Risk-averse

9.7 Reference

Larroyd, P.V., Matos, V.L., Diniz, A.L., & Borges, C.L.T. (2022). “Tackling the Seasonal and Stochastic Components in Hydro-Dominated Power Systems with High Renewable Penetration.” *Energies*, 15(3), 1115. <https://doi.org/10.3390/en15031115>

Part IV: Cut Management and Convergence

10. Cut Generation and Aggregation

10.1 Dual Variable Extraction

After solving the stage t subproblem for state $\hat{x}_{t-1}$ and scenario ω_t , extract dual variables from the optimal LP solution:

Constraint	Dual Variable	Notation	Units
Water balance (hydro h)	π_h^{wb}	Shadow price of storage	\$/hm ³
AR lag fixing (hydro h , lag ℓ)	$\pi_{h,\ell}^{lag}$	Shadow price of inflow lag	\$/ (m ³ /s)
Generic constraint (constraint c)	π_c^{gen}	Shadow price of generic constraint	depends

Sign Convention: For minimization LPs with $\leq$ constraints, $\pi \geq 0$. For equality constraints (water balance), the sign depends on constraint orientation. See Section 3.4 for detailed sign convention for cut coefficient computation.

10.2 Cut Coefficient Computation

From the dual variables, compute cut coefficients for state variables:

Storage coefficient (marginal value of water):

$$\beta_{t,h}^v = \pi_h^{wb} \cdot \zeta$$

Unit Analysis: - π_h^{wb} has units $\$/\text{hm}^3$ (shadow price of the water balance constraint, which is in hm^3)
 - ζ has units $\text{hm}^3/(\text{m}^3/\text{s})$ (converts flow rate to volume over the stage) - Product $\beta_{t,h}^v$ has units $\$/(\text{m}^3/\text{s})$
 — but this requires clarification:

Why does ζ appear? The water balance constraint is written in terms of storage (v_h , in hm^3), but the incoming flow state variables ($\hat{a}_{h,\ell}$) are in m^3/s . The factor ζ ensures dimensional consistency when the cut coefficient is applied to the state variable.

Alternative interpretation: Some formulations write the water balance in terms of total inflow volume ($A_h = a_h \cdot \zeta$), in which case the dual π_h^{wb} is already in $\$/\text{hm}^3$ and no ζ factor is needed. POWE.RS uses the flow-rate formulation for consistency with PAR(p) model parameters.

AR lag coefficient (marginal value of historical inflow information):

$$\beta_{t,h,\ell}^{lag} = \pi_{h,\ell}^{lag}$$

Cut intercept (computed to make the cut pass through the trial point):

$$\alpha_t = Q_t(\hat{x}_{t-1}, \omega_t) - \sum_{h \in \mathcal{H}} \beta_{t,h}^v \cdot \hat{v}_h - \sum_{h \in \mathcal{H}} \sum_{\ell=1}^{p_h} \beta_{t,h,\ell}^{lag} \cdot \hat{a}_{h,\ell}$$

where $Q_t(\hat{x}_{t-1}, \omega_t)$ is the optimal objective value of the stage t subproblem.

10.3 Single-Cut Aggregation

For **risk-neutral** problems with single-cut aggregation, compute expectation over scenarios:

Aggregated intercept:

$$\bar{\alpha}_{t-1} = \sum_{\omega \in \Omega_t} p(\omega) \cdot \alpha_t(\omega)$$

Aggregated storage coefficients:

$$\bar{\beta}_{t-1,h}^v = \sum_{\omega \in \Omega_t} p(\omega) \cdot \beta_{t,h}^v(\omega)$$

Aggregated lag coefficients:

$$\bar{\beta}_{t-1,h,\ell}^{lag} = \sum_{\omega \in \Omega_t} p(\omega) \cdot \beta_{t,h,\ell}^{lag}(\omega)$$

where $p(\omega)$ is the probability of scenario ω .

10.4 Multi-Cut Formulation (DEFERRED)

The multi-cut variant creates one cut per scenario instead of aggregating:

$$\theta_{t-1,\omega} \geq \beta_{t-1 \rightarrow t} \cdot \left(\alpha_t(\omega) + \sum_h \beta_{t,h}^v(\omega) \cdot v_h + \sum_{h,\ell} \beta_{t,h,\ell}^{lag}(\omega) \cdot a_{h,\ell} \right) \quad \forall \omega \in \Omega_t$$

with aggregation in the objective:

$$\theta_{t-1} = \sum_{\omega \in \Omega_t} p(\omega) \cdot \theta_{t-1,\omega}$$

Trade-offs: - More cuts per iteration (one per scenario vs. one total) - Potentially faster convergence in early iterations - Higher LP solve time due to more constraints

See [Appendix C](#) for planned implementation details.

10.5 Cut Addition Algorithm

Algorithm: Backward Pass Cut Generation

1. For each stage t from $T - 1$ down to 1:
 - For each trial point $\hat{x}$ from forward pass:
 - Initialize `cuts_for_this_point` = []
 - For each scenario $\omega \in \Omega_t$:
 - * Solve subproblem: $(Q_\omega, \pi_\omega) = \text{solve_subproblem}(t, \hat{x}, \omega)$
 - * Compute per-scenario cut coefficients:

$$\beta_\omega^v = \text{compute_storage_coefficients}(\pi_\omega)$$

$$\beta_\omega^{lag} = \text{compute_lag_coefficients}(\pi_\omega)$$

$$\alpha_\omega = Q_\omega - \beta_\omega^v \cdot \hat{x}.\text{storage} - \beta_\omega^{lag} \cdot \hat{x}.\text{lags}$$
 - * Append to cuts: `cuts_for_this_point.append(( $\alpha_\omega, \beta_\omega^v, \beta_\omega^{lag}$ ))`
 - Aggregate (single-cut formulation):

$$(\bar{\alpha}, \bar{\beta}^v, \bar{\beta}^{lag}) = \text{aggregate_cuts}(\text{cuts_for_this_point}, \text{probabilities})$$
 - Add to stage $t - 1$:

$$\text{add_cut_to_stage}(t - 1, \text{intercept} = \bar{\alpha}, \text{coefs} = (\bar{\beta}^v, \bar{\beta}^{lag}))$$

10.7 Cut Validity

A cut is **valid** if it provides a lower bound on the true cost-to-go function:

$$\alpha_k + \beta_k^\top x \leq V_{t+1}(x) \quad \forall x \in \mathcal{X}_t$$

Theorem: The cuts generated by SDDP are valid under: 1. Convexity of stage subproblems 2. Relatively complete recourse (feasibility for all scenarios) 3. Correct dual extraction from optimal LP solutions

11. Cut Selection Strategies

11.1 Motivation

As SDDP iterations progress, the number of Benders cuts grows linearly ($\mathcal{O}(\text{iterations} \times \text{forward_passes})$). Many cuts become redundant (dominated by newer, tighter cuts). Cut selection removes inactive cuts to: 1. Reduce LP solve time (fewer constraints) 2. Improve numerical stability (remove near-parallel constraints) 3. Maintain memory efficiency

11.2 Cut Activity Definition

A cut k at stage t is **active** at state $\hat{x}$ if it is binding at the optimal solution:

$$\theta^* = \alpha_k + \beta_k^\top \hat{x}$$

Equivalently, the cut constraint has **positive dual multiplier** $\lambda_k > 0$.

A cut is **dominated** if there exists no visited state where it is active.

11.3 Level-1 Cut Selection

Definition: A cut is **Level-1** if it was active at least once during the entire algorithm execution.

Algorithm: Level-1 Cut Selection

1. After each backward pass:
 - For each stage t :
 - For each cut k in stage t :
 - * If cut k was active in current iteration:
 - Mark cut k as “used”
 - 2. Periodically (every N iterations):
 - For each stage t :
 - For each cut k in stage t :
 - * If cut k was never “used”:
 - Deactivate cut k (set bound to $-\infty$)

Properties: - Simple to implement (just track “ever active” flag) - Preserves convergence guarantee - May retain some dominated cuts (active once, never again)

11.4 Limited Memory Level-1 (LML1)

Definition: For each visited state, keep only the **most recently active** cut.

Algorithm: Limited Memory Level-1

1. After backward pass at iteration i :
 - For each visited state $\hat{x}$ in forward pass:
 - For each stage t :
 - * Identify which cut k^* was active at $\hat{x}$
 - * Mark k^* with timestamp i
 - 2. Periodically:
 - For each stage t :
 - For each cut k :
 - * If $k.\text{timestamp} < \text{current_iteration} - \text{memory_window}$:
 - Deactivate cut k

Properties: - More aggressive than Level-1 - Memory window controls retention period - Still preserves finite convergence (with probability 1)

11.5 Dominated Cut Detection

A cut k is **dominated** by a set of cuts $\mathcal{S}$ if:

$$\alpha_k + \beta_k^\top x \leq \max_{j \in \mathcal{S}} \{ \alpha_j + \beta_j^\top x \} \quad \forall x \in \mathcal{X}$$

Practical Detection (at visited states):

For each visited state $\hat{x}$, compute:

$$\Delta_k(\hat{x}) = \max_{j \neq k} \{\alpha_j + \beta_j^\top \hat{x}\} - (\alpha_k + \beta_k^\top \hat{x})$$

If $\Delta_k(\hat{x}) > \epsilon$ for all visited states, cut k is **dominated**.

Algorithm: Dominated Cut Detection

1. For each stage t :
 - Let `states = visited_states[t]`
 - For each cut k :
 - Set `dominated = true`
 - For each state $\hat{x}$ in `states`:
 - * Compute: $\text{value}_k = \alpha_k + \beta_k^\top \hat{x}$
 - * Compute: $\text{max_other} = \max_{j \neq k} (\alpha_j + \beta_j^\top \hat{x})$
 - * If $\text{value}_k \geq \text{max_other} - \text{threshold}$:
 - Set `dominated = false`
 - Break
 - If `dominated`:
 - * Deactivate cut k

11.6 Threshold Parameter

The `threshold` parameter in cut selection controls the minimum violation to consider a cut active:

$$\text{cut } k \text{ is active at } \hat{x} \iff \theta^* - (\alpha_k + \beta_k^\top \hat{x}) < \text{threshold}$$

Threshold Value	Behavior
0	Only strictly binding cuts are active
1e-6	Near-binding cuts included (numerical tolerance)
1e-3	Moderately loose cuts included
Large	All cuts considered active (no selection)

Recommended: `threshold = 0` with numerical tolerance handled separately.

11.7 Cut Selection Configuration

```
{
  "training": {
    "cut_selection": {
      "enabled": true,
      "method": "domination",
      "threshold": 0,
      "check_frequency": 10,
      "memory_window": 50
    }
  }
}
```

Parameter	Description
<code>method</code>	"level1", "lml1", or "domination"
<code>threshold</code>	Minimum violation for activity
<code>check_frequency</code>	Iterations between cut selection runs
<code>memory_window</code>	For LML1: iterations to retain inactive cuts

11.8 Convergence Guarantee

Theorem (Guigues & Bandarra, 2019): Under Level-1 or LML1 cut selection, SDDP with finitely many scenarios converges to the optimal value function with probability 1.

Key insight: Removing cuts that are never active at any visited state does not affect the quality of the outer approximation at those states. As the set of visited states becomes dense, the approximation converges.

11.9 Reference

Guigues, V., & Bandarra, M.P. (2019). “Single cut and multicut SDDP with cut selection for multistage stochastic linear programs: convergence proof and numerical experiments.” *arXiv:1902.06757*.
<https://arxiv.org/abs/1902.06757>

12. Stopping Rules Evaluation

12.1 Available Stopping Rules

SDDP can terminate based on multiple criteria. Each rule is evaluated independently, and the `stopping_mode` determines how they combine:

- "any": Stop when **any** rule triggers (OR logic)
- "all": Stop when **all** rules trigger (AND logic)

12.2 Iteration Limit (Mandatory)

Configuration:

```
{"type": "iteration_limit", "limit": 50}
```

Evaluation:

$$\text{STOP} \iff k \geq k_{max}$$

where k is the current iteration and k_{max} is the limit.

Purpose: Safety bound to prevent infinite loops. **Must always be included.**

12.3 Time Limit

Configuration:

```
{"type": "time_limit", "seconds": 3600}
```

Evaluation:

$$\text{STOP} \iff t_{elapsed} \geq t_{max}$$

Implementation: Check wall-clock time at end of each iteration.

12.4 Statistical Stopping

Configuration:

```
{  
  "type": "statistical",  
  "num_replications": 100,  
  "iteration_period": 5,  
  "z_score": 1.96  
}
```

Algorithm:

1. Every `iteration_period` iterations, run `num_replications` Monte Carlo simulations using the current policy

2. Compute sample statistics:

$$\bar{z} = \frac{1}{M} \sum_{m=1}^M z^{(m)}$$

$$s_z = \sqrt{\frac{1}{M-1} \sum_{m=1}^M (z^{(m)} - \bar{z})^2}$$

where $z^{(m)}$ is the total cost of simulation m .

3. Compute confidence interval half-width:

$$w = z_{\alpha/2} \cdot \frac{s_z}{\sqrt{M}}$$

where $z_{\alpha/2}$ is the z-score (1.96 for 95% confidence).

4. **Stopping condition** (for minimization):

$$\text{STOP} \iff \underline{z}^k \geq \bar{z} - w$$

i.e., the deterministic lower bound lies within the confidence interval of the simulated upper bound.

Interpretation: With confidence $1 - \alpha$, the optimal value lies in $[\underline{z}^k, \bar{z} + w]$. If $\underline{z}^k \geq \bar{z} - w$, the gap is statistically zero.

Warning: Not Recommended for Production Use

The statistical stopping rule has several serious limitations:

1. **Half-width depends on replications:** With more replications M , the half-width $w \propto 1/\sqrt{M}$ shrinks, making the stopping condition easier to satisfy. This creates a perverse incentive: fewer replications = faster “convergence.”
2. **Non-normality:** Cost distributions in stochastic optimization are often heavy-tailed or multimodal, violating the normality assumption underlying the z-score confidence interval.
3. **Sequential testing inflation:** Testing the stopping condition repeatedly at each period inflates the Type I error rate. The nominal 95% confidence level does not hold under repeated testing.
4. **Risk-averse incompatibility:** For risk-averse problems, the lower bound is not valid (see Section 9.12), making this rule fundamentally flawed.

Recommended Alternatives: - Use **bound_stalling** (Section 6.5) for deterministic convergence monitoring - Use **simulation** stopping (Section 6.6) with explicit gap tolerance - For risk-averse problems, rely on iteration limits combined with policy stability metrics

12.5 Bound Stalling

Configuration:

```
{
  "type": "bound_stalling",
  "iterations": 10,
  "tolerance": 0.0001
}
```

Evaluation:

Track the deterministic lower bound $\underline{z}^k$ over iterations. Compute relative improvement:

$$\Delta_k = \frac{\underline{z}^k - \underline{z}^{k-\tau}}{\max(1, |\underline{z}^k|)}$$

where τ is the `iterations` parameter.

Stopping condition:

$$\text{STOP} \iff |\Delta_k| < \text{tolerance for consecutive } \tau \text{ iterations}$$

Interpretation: The bound has plateaued—further iterations provide diminishing returns.

12.6 Simulation-Based Stopping (Recommended)

Configuration:

```
{
  "type": "simulation",
  "replications": 100,
  "period": 20,
  "distance_tol": 0.01,
  "bound_tol": 0.0001
}
```

Algorithm:

1. **Check bound stability first:**

$$\text{Bound stable} \iff |\underline{z}^k - \underline{z}^{k-5}| < \text{bound_tol} \times \max(1, |\underline{z}^k|)$$

2. **If bound is stable**, run simulation and compare to previous simulation:

Let $\mathbf{x}^{new}$ and $\mathbf{x}^{old}$ be vectors of simulated state trajectories (or costs per stage).

Compute normalized distance:

$$d = \sqrt{\sum_i \left(\frac{x_i^{new} - x_i^{old}}{\max(1, |x_i^{old}|)} \right)^2}$$

3. **Stopping condition:**

$$\text{STOP} \iff \text{Bound stable} \wedge d < \text{distance_tol}$$

Interpretation: Both the outer approximation (bound) and the policy (simulated decisions) have stabilized.

Why recommended: Combines theoretical convergence indicator (bound) with practical policy quality (simulation), avoiding premature termination from statistical noise.

12.7 Combining Rules

Mode: "any" (default):

$$\text{STOP} \iff \text{Rule}_1 \vee \text{Rule}_2 \vee \dots$$

First rule to trigger causes termination.

Mode: "all":

$$\text{STOP} \iff \text{Rule}_1 \wedge \text{Rule}_2 \wedge \dots$$

All rules must trigger simultaneously.

Example (conservative setup):

```
{
  "stopping_rules": [
    {"type": "iteration_limit", "limit": 500},
    {"type": "simulation", "replications": 100, "period": 20, "distance_tol": 0.01, "bound_tol": 0.0001}
  ],
  "stopping_mode": "any"
}
```

This runs until simulation-based convergence OR 500 iterations, whichever comes first.

12.8 Output on Termination

When any stopping rule triggers, the output includes:

Field	Description
stopping_rule	Which rule triggered
final_iteration	Iteration count at termination
lower_bound	Final deterministic lower bound
upper_bound	Final simulated upper bound (if available)
gap_percent	$(\bar{z} - \underline{z})/ \bar{z} \times 100$

Part V: Advanced Formulations

13. Discount Rate Formulation

13.1 Motivation

The discount rate $\beta \in (0, 1]$ captures the time value of money or risk preference, where future costs are valued less than present costs. This is essential for: 1. **Infinite horizon problems**: Ensuring convergence of the value function 2. **Economic consistency**: Reflecting opportunity cost of capital 3. **Risk adjustment**: Implicitly reducing weight of distant uncertain outcomes

13.2 Discounted Bellman Equation

The standard risk-neutral Bellman recursion with discount factor β is:

$$V_t(x_{t-1}) = \mathbb{E}_{\omega_t} \left[\min_{x_t \in \mathcal{X}_t(\omega_t)} \{c_t(x_t, u_t) + \beta \cdot V_{t+1}(x_t)\} \right]$$

where: - $c_t(x_t, u_t)$ is the immediate cost at stage t - $V_{t+1}(x_t)$ is the future cost function (cost-to-go) - $\beta = \frac{1}{1+r}$ with r being the discount rate per stage

Formulation Note: The discount factor β multiplies only the **future cost** V_{t+1} , not the immediate cost c_t . This is the standard SDDP convention:

- **Immediate cost** c_t : Not discounted (incurred “now” at stage t)
- **Future cost** $\beta \cdot V_{t+1}$: Discounted to present value at stage t

This is mathematically equivalent to computing all costs at “time 0” (stage 1) present value, where stage t costs are multiplied by $\prod_{s=1}^{t-1} \beta_s$ in the objective. The Bellman formulation above is the recursive form that SDDP exploits.

Alternative formulations (not used in POWE.RS): - Some formulations discount both immediate and future cost by β_t within the expectation - The choice affects cut coefficient scaling but not the optimal policy

13.3 Stage-Dependent Discount Rates

In POWE.RS, discount rates are specified per **transition** in `stages.json`:

```
{  
  "transitions": [  
    {"source_id": 0, "target_id": 1, "probability": 1.0, "discount_rate": 0.005},  
    {"source_id": 1, "target_id": 2, "probability": 1.0, "discount_rate": 0.005}  ]  
}
```

]

}

The discount factor for transition from stage t to stage $t + 1$ is:

$$\beta_{t \rightarrow t+1} = \frac{1}{1 + r_{t \rightarrow t+1}}$$

13.4 Modified Stage Subproblem

The stage t subproblem with discounting becomes:

$$Q_t(x_{t-1}, \omega_t) = \min_{x_t, u_t, \theta} \{c_t(x_t, u_t) + \theta\}$$

subject to: - All standard constraints (load balance, hydro balance, etc.) - **Discounted Benders cuts**:

$$\theta \geq \beta_{t \rightarrow t+1} \cdot \left(\alpha_i + \sum_h \beta_{i,h}^v \cdot v_h + \sum_{h,\ell} \beta_{i,h,\ell}^{lag} \cdot a_{h,\ell} \right) \quad \forall i$$

13.5 Cumulative Discounting

For a path from stage 1 to stage T , the cumulative discount factor is:

$$\beta_{1 \rightarrow T} = \prod_{t=1}^{T-1} \beta_{t \rightarrow t+1}$$

The present value at stage 1 of costs incurred at stage T is:

$$\text{PV}_1[c_T] = \beta_{1 \rightarrow T} \cdot c_T$$

13.6 Lower Bound Computation with Discounting

The deterministic lower bound at iteration k is computed as:

$$\underline{z}^k = c_1(\hat{x}_1^k) + \theta_1^k$$

where θ_1^k is the optimal value of the future cost variable at stage 1, which already includes all discounting through the cut coefficients.

13.7 Upper Bound (Simulation) with Discounting

When simulating the policy to estimate the upper bound:

$$\bar{z}^k = \frac{1}{M} \sum_{m=1}^M \sum_{t=1}^T \beta_{1 \rightarrow t} \cdot c_t(\hat{x}_t^{k,m})$$

where $\beta_{1 \rightarrow 1} = 1$ and $\beta_{1 \rightarrow t} = \prod_{s=1}^{t-1} \beta_{s \rightarrow s+1}$.

13.8 Implementation Notes

- **Cut storage:** Cuts are stored in **undiscounted** form. Discounting is applied when adding to LP.
- **Cut coefficients in LP:** The LP stores $\beta \cdot \alpha$ and $\beta \cdot \beta^v$, not the raw values.
- **Consistent units:** All bounds and gaps are reported in present value terms (stage 1 currency).

Bound Interpretation: Both the lower bound $\underline{z}$ and upper bound $\bar{z}$ represent total expected cost expressed in **present value at stage 1**. This means: - A lower bound of \$100M means the optimal policy costs at least \$100M in stage-1 dollars - Future costs are already discounted: a \$1M cost at stage 12 with $\beta_{1 \rightarrow 12} = 0.95$ contributes \$0.95M to the bounds - Comparisons between bounds and between iterations are valid because they use consistent discounting - When reporting per-stage costs in simulation outputs, POWERS reports both **nominal** (undiscounted) and **present value** costs

14. Infinite Periodic Horizon Formulation

14.1 Motivation

Standard finite-horizon SDDP has a terminal condition $V_{T+1}(x) = 0$, causing “end-of-world” effects where the algorithm empties reservoirs toward the horizon. For long-term planning, an **infinite periodic horizon** better represents the ongoing nature of hydrothermal operations.

14.2 Periodic Structure

Consider a system with 12 monthly stages that repeat annually. Let $\tau(t)$ denote the **season** (position within the cycle) for stage t :

$$\tau(t) = (t - 1) \mod 12 + 1 \in \{1, 2, \dots, 12\}$$

Stages with the same season share structural properties (demand patterns, inflow statistics).

14.3 Cycle Detection

The algorithm detects cycles by analyzing the transition graph in **stages.json**. A **cycle** exists when a transition points to a stage with ID less than or equal to the source (backward edge in a DAG sense).

Example:

```
{
  "transitions": [
    {"source_id": 0, "target_id": 1, "probability": 1.0},
    ...
    {"source_id": 59, "target_id": 48, "probability": 1.0, "discount_rate": 0.05}
  ]
}
```

Here, stage 59 transitions back to stage 48, creating a 12-stage cycle.

14.4 Discounting for Convergence

For convergence, the cycle must include a **discount factor** $\beta < 1$ on the return edge. This ensures:

$$\lim_{n \rightarrow \infty} \beta^n V_t(x) = 0$$

The cumulative discount around one full cycle must satisfy:

$$\beta_{cycle} = \prod_{t \in cycle} \beta_{t \rightarrow t+1} < 1$$

Typical setup: Monthly discount rate of 0.5% gives $\beta = 1/1.005 \approx 0.995$, and annual discount $\beta_{cycle} = 0.995^{12} \approx 0.94$.

14.5 Cut Sharing Within Cycles

Stages in the same position of the cycle share their value function approximation. Let $\mathcal{C}_\tau = \{t : \tau(t) = \tau\}$ be all stages with season τ .

Cut sharing rule: A cut generated at stage $t \in \mathcal{C}_\tau$ is applicable to all stages in $\mathcal{C}_\tau$:

$$\underline{V}_\tau(x) = \max_{k \in \mathcal{K}_\tau} \{\alpha_k + \beta_k^\top x\}$$

Implementation: The cut pool is indexed by season $\tau \in \{1, \dots, 12\}$, not by absolute stage ID.

14.6 Fixed-Point Iteration

The infinite-horizon SDDP finds the fixed point of the Bellman operator:

$$V_\tau = T_\tau V_{\tau+1}$$

where T_τ is the one-stage Bellman operator for season τ :

$$(T_\tau V)(x) = \mathbb{E}_{\omega_\tau} \left[\min_{x'} \{c_\tau(x', u) + \beta \cdot V(x')\} \right]$$

Convergence criterion: Value functions have converged when:

$$\|V_\tau^{k+1} - V_\tau^k\|_\infty < \delta_{cycle}$$

for all seasons τ , where δ_{cycle} is the `cycle_discretization_delta` tolerance.

14.7 Modified Forward Pass

In infinite horizon, the forward pass continues until the discounted contribution becomes negligible:

Algorithm: Infinite Horizon Forward Pass

1. For forward pass m :
 - Initialize: $t = 0$, `cumulative_discount` = 1.0, $x = x_0$, `total_cost` = 0
 - While `cumulative_discount` > `tolerance` AND $t < \text{max_horizon_length}$:
 - Sample ω_t from stage $\tau(t)$
 - Solve subproblem: $(x', \text{cost}) = \text{solve_subproblem}(t, x, \omega_t)$
 - Update: `total_cost` += `cumulative_discount` * `cost`
 - Update: $x = x'$
 - Increment: $t = t + 1$
 - Update discount: `cumulative_discount` *= $\beta_{t-1 \rightarrow t}$

The `max_horizon_length` provides a safety bound (e.g., 240 stages = 20 years for monthly).

14.8 Backward Pass Modifications

Stopping condition: The backward pass stops when it completes a full cycle with no significant improvement:

$$\max_\tau |z^{k,\tau} - z^{k-12,\tau}| < \delta_{cycle}$$

Cut generation: Same as finite horizon, but cuts are added to the season's cut pool, not a specific stage.

14.9 Configuration

```
{
  "horizon": {
    "mode": "infinite_periodic",
    "max_horizon_length": 240,
    "cycle_discretization_delta": 0.1
  }
}
```

Parameter	Description
mode	"infinite_periodic" enables this formulation
max_horizon_length	Maximum stages in forward pass
cycle_discretization_delta	Convergence tolerance for cycle

14.10 Reference

Costa, B.S., de Matos, V.L., Philpott, A.B. (2025). "SDDP.jl approaches for infinite horizon problems." *Trends in Computational and Applied Mathematics*, 11(1). <https://doi.org/10.5540/03.2025.011.01.0355>

15. Upper Bound Evaluation LP (Inner Approximation / SIDP)

15.1 Motivation

Standard SDDP provides only a **lower bound** (outer approximation) through cuts. For convergence verification, we need an **upper bound** (inner approximation). This is especially important for:

1. **Risk-averse problems:** CVaR objectives cannot be estimated via Monte Carlo
2. **Convergence certificates:** Gap = $\bar{z} - \underline{z}$ provides true optimality measure
3. **Conservative policies:** Inner approximation gives "at most Y" guarantees

15.2 Vertex-Based Inner Approximation

The inner approximation $\bar{V}_t(x)$ is constructed from **vertices** (visited state-value pairs):

$$\mathcal{V}_t = \{(x^{(1)}, \bar{v}^{(1)}), (x^{(2)}, \bar{v}^{(2)}), \dots, (x^{(n)}, \bar{v}^{(n)})\}$$

where each vertex stores: - $x^{(i)}$: State vector visited during forward passes - $\bar{v}^{(i)}$: Upper bound on cost-to-go from that state (computed recursively)

15.3 Lipschitz Interpolation

For a new state x not in $\mathcal{V}_t$, the upper bound is computed via Lipschitz interpolation:

$$\bar{V}_t(x) = \min_{(x^{(i)}, \bar{v}^{(i)}) \in \mathcal{V}_t} \{\bar{v}^{(i)} + L_t \cdot \|x - x^{(i)}\|_1\}$$

where L_t is the Lipschitz constant for stage t .

Interpretation: The upper bound at x is the minimum over all vertices of "vertex value plus distance penalty."

15.4 Lipschitz Constant Computation

The Lipschitz constant bounds the maximum rate of change of the value function. For SDDP with penalty-based feasibility:

Backward accumulation:

$$L_T = c_{max}^{penalty}$$

$$L_t = L_{t+1} + c_{max}^{penalty,t}$$

where $c_{max}^{penalty,t}$ is the maximum penalty coefficient at stage t (e.g., deficit penalty in \$/MWh).

Example: With deficit penalty 1000 \$/MWh over 5 stages: - $L_5 = 1000$ - $L_4 = 2000$ - $L_3 = 3000$ - $L_2 = 4000$ - $L_1 = 5000$

15.5 Vertex Value Computation

During upper bound evaluation (backward pass variant):

At terminal stage T :

$$\bar{v}^{(i)} = c_T(x^{(i)}) \quad (\text{immediate cost only})$$

At stage $t < T$:

1. For each vertex $(x^{(i)}, \cdot) \in \mathcal{V}_t$:
2. Solve the stage subproblem with $x = x^{(i)}$
3. Compute future cost using inner approximation at stage $t + 1$:

$$\bar{\theta} = \bar{V}_{t+1}(x_{t+1}^*)$$

4. Set vertex value:

$$\bar{v}^{(i)} = c_t(x^{(i)}) + \beta \cdot \bar{\theta}$$

15.6 Upper Bound Evaluation LP

For policy simulation with inner approximation, the stage LP is modified:

Standard LP (outer approximation):

$$\begin{aligned} & \min c_t^\top x_t + \theta \\ \text{s.t. } & \theta \geq \alpha_k + \beta_k^\top x_t \quad \forall k \text{ (cuts)} \end{aligned}$$

Inner approximation LP:

$$\begin{aligned} & \min c_t^\top x_t + \bar{\theta} \\ \text{s.t. } & \bar{\theta} \leq \bar{v}^{(i)} + L_t \sum_j |x_{t,j} - x_j^{(i)}| \quad \forall i \text{ (vertices)} \end{aligned}$$

The absolute value can be linearized using standard techniques:

$$\begin{aligned} |x_j - x_j^{(i)}| &= u_j^{(i)+} + u_j^{(i)-} \\ x_j - x_j^{(i)} &= u_j^{(i)+} - u_j^{(i)-} \\ u_j^{(i)+}, u_j^{(i)-} &\geq 0 \end{aligned}$$

15.7 Linearized Upper Bound LP

Additional Variables (per vertex i , per state component j):

Variable	Domain	Description
$u_j^{(i)+}$	≥ 0	Positive deviation from vertex i in dimension j
$u_j^{(i)-}$	≥ 0	Negative deviation from vertex i in dimension j
$\bar{\theta}$	free	Upper bound on future cost

Constraints:

For each vertex $i \in \mathcal{V}_t$:

$$\bar{\theta} \leq \bar{v}^{(i)} + L_t \sum_j (u_j^{(i)+} + u_j^{(i)-})$$

$$x_j - x_j^{(i)} = u_j^{(i)+} - u_j^{(i)-} \quad \forall j$$

15.8 Gap Computation

At each iteration k (when upper bound is evaluated):

Lower bound (from cuts at stage 1):

$$\underline{z}^k = c_1(\hat{x}_1) + \underline{V}_2(\hat{x}_1)$$

Upper bound (from vertices):

$$\bar{z}^k = c_1(\hat{x}_1) + \bar{V}_2(\hat{x}_1)$$

Gap:

$$\text{gap}^k = \frac{\bar{z}^k - \underline{z}^k}{\max(1, |\bar{z}^k|)} \times 100\%$$

Convergence: As $k \rightarrow \infty$, $\text{gap}^k \rightarrow 0$ for convex problems with finitely many scenarios.

15.9 Vertex Storage

Vertices are stored in `policy/vertices/stage_XXX.bin` with the following schema:

Field	Type	Description
<code>state</code>	<code>[f64; n_state]</code>	State vector at vertex
<code>value</code>	<code>f64</code>	Upper bound value at vertex
<code>iteration</code>	<code>u32</code>	Iteration when vertex was created
<code>lipschitz</code>	<code>f64</code>	Per-vertex Lipschitz constant (if variable)

15.10 Configuration

```
{
  "upper_bound_evaluation": {
    "enabled": true,
    "initial_iteration": 10,
    "interval_iterations": 5,
    "lipschitz": {
```

```

    "mode": "auto",
    "fallback_value": 10000.0,
    "scale_factor": 1.1
  }
}
}

```

15.11 Computational Considerations

Aspect	Impact
Vertices per stage	Typically $\mathcal{O}(\text{iterations} \times \text{forward_passes})$
LP size increase	$2 \times n_{\text{state}} \times n_{\text{vertices}}$ additional variables
Evaluation frequency	Trade-off between gap accuracy and runtime
Memory	Vertices stored separately from cuts

Recommendation: Enable upper bound evaluation every 5-10 iterations after initial burn-in period (10+ iterations) for convergence monitoring without excessive overhead.

15.12 References

Costa, B.S., & Leclère, V. (2023). “Lipschitz-based Inner Approximation of Risk Measures.” *Optimization Online*. <https://optimization-online.org/?p=23738>

Philpott, A.B., de Matos, V.L., & Finardi, E.C. (2013). “On solving multistage stochastic programs with coherent risk measures.” *Operations Research*, 61(4), 957-970. <https://doi.org/10.1287/opre.2013.1200>

16. Risk-Averse SDDP (CVaR) Formulation

16.1 Motivation

Risk-neutral SDDP minimizes expected cost, which can lead to policies that perform poorly in adverse scenarios. **Risk-averse SDDP** incorporates a coherent risk measure (typically CVaR) to protect against tail risks.

16.2 Conditional Value-at-Risk (CVaR)

For a random variable Z representing cost and confidence level $\alpha \in (0, 1]$:

$$\text{CVaR}_\alpha(Z) = \min_{\eta \in \mathbb{R}} \left\{ \eta + \frac{1}{\alpha} \mathbb{E}[(Z - \eta)^+] \right\}$$

where $(Z - \eta)^+ = \max(0, Z - \eta)$ captures the excess cost above threshold η .

Interpretation: CVaR_α is the expected cost in the worst α -fraction of scenarios.

α	Risk Posture	Meaning
1.0	Risk-neutral	$\text{CVaR}_1 = \mathbb{E}[Z]$ (expected value)
0.5	Moderately risk-averse	Average of worst 50% of outcomes
0.2	Risk-averse	Average of worst 20% of outcomes
0.05	Highly risk-averse	Average of worst 5% of outcomes

16.3 Convex Combination Risk Measure (SDDP.jl Convention)

SDDP.jl uses a convex combination of expectation and CVaR:

$$\rho^{\lambda, \alpha}[Z] = (1 - \lambda)\mathbb{E}[Z] + \lambda \cdot \text{CVaR}_\alpha[Z]$$

where: - $\lambda \in [0, 1]$: Risk aversion weight (0 = risk-neutral, 1 = pure CVaR) - $\alpha \in (0, 1]$: CVaR confidence level

16.4 Dual Representation of Convex Risk Measures

Convex risk measures have a **dual representation** that is essential for computing risk-averse cuts:

$$\mathbb{F}[Z] = \sup_{q \in \mathcal{M}(p)} \mathbb{E}_q[Z] - \alpha(p, q)$$

where: - $\mathcal{M}(p) \subseteq \mathcal{P}$ is a convex subset of the probability simplex - $\alpha(p, q)$ is a concave penalty function - $\mathcal{P} = \{p \geq 0 : \sum_\omega p_\omega = 1\}$

Interpretation: The dual computes the expectation with respect to the **worst** probability vector q within the set $\mathcal{M}$, less a penalty term $\alpha(p, q)$.

CVaR Dual Representation For CVaR_α , the dual representation is:

$$\text{CVaR}_\alpha[Z] = \sup_{\mu \in \mathcal{M}_\alpha(p)} \mathbb{E}_\mu[Z]$$

where the risk set $\mathcal{M}_\alpha(p)$ is:

$$\mathcal{M}_\alpha(p) = \left\{ \mu \geq 0 : \sum_\omega \mu_\omega = 1, \mu_\omega \leq \frac{p_\omega}{\alpha} \forall \omega \right\}$$

Notation Convention: We use μ (Greek mu) for the risk-adjusted probability measure to avoid confusion with turbined flow q .

The penalty $\alpha(p, \mu) = 0$ for CVaR (no penalty term).

Interpretation: CVaR puts more probability weight on the worst outcomes, with each scenario receiving at most p_ω/α probability mass. For small α , only the worst scenarios receive significant weight.

EAVaR Dual Representation For the convex combination $\rho^{\lambda, \alpha}[Z] = (1 - \lambda)\mathbb{E}[Z] + \lambda \cdot \text{CVaR}_\alpha[Z]$:

$$\mathcal{M}^{EAVaR}(p) = \left\{ \mu \geq 0 : \sum_\omega \mu_\omega = 1, \mu_\omega \leq (1 - \lambda)p_\omega + \frac{\lambda p_\omega}{\alpha} \forall \omega \right\}$$

16.5 Risk-Averse Subgradient Theorem

The key theorem for computing risk-averse cuts:

Theorem (Risk-Averse Subgradient): Let $V(x, \omega)$ be convex with respect to x for all fixed $\omega \in \Omega$, and let $\lambda(\tilde{x}, \omega)$ be a subgradient of $V(x, \omega)$ at $x = \tilde{x}$.

If $\mu^* = \arg\max_{\mu \in \mathcal{M}(p)} \mathbb{E}_\mu[V(\tilde{x}, \omega)] - \alpha(p, \mu)$, then:

$$\sum_{\omega \in \Omega} \mu_\omega^* \cdot \lambda(\tilde{x}, \omega)$$

is a subgradient of $\mathbb{F}[V(x, \omega)]$ at $\tilde{x}$.

Application to Cut Generation:

In SDDP, the subgradients $\lambda(\tilde{x}, \omega)$ are the cut coefficients $\beta_t(\omega)$ obtained from LP duals. The risk-averse cut coefficients are computed as:

$$\bar{\beta}_{t-1,h} = \sum_{\omega \in \Omega_t} \mu_\omega^* \cdot \beta_{t,h}(\omega)$$

where μ^* is the optimal dual probability vector computed from the scenario costs $\{Q_t(\hat{x}, \omega)\}_{\omega \in \Omega_t}$.

16.6 Risk-Averse Bellman Equation

The risk-averse value function satisfies:

$$V_t(x_{t-1}) = \rho^{\lambda, \alpha} \left[\min_{x_t} \{c_t^\top x_t + V_{t+1}(x_t) : (x_t, x_{t-1}) \text{ feasible}\} \right]$$

16.7 Cut Generation with Risk Measures

For each visited state $\hat{x}_{t-1}$, compute the risk-averse cut as follows:

Step 1: Solve subproblems for all realizations $\omega \in \Omega_t$:

$$Q_t(\hat{x}_{t-1}, \omega) = \min_{x_t} \{c_t^\top x_t + \theta_t : \text{constraints}\}$$

Extract dual solutions $\pi_t(\omega)$ and compute per-scenario cut coefficients: - Intercept: $\alpha_t(\omega) = Q_t(\hat{x}_{t-1}, \omega) - \beta_t(\omega)^\top \hat{x}_{t-1}$
- Coefficients: $\beta_t(\omega) = -\pi_t(\omega)$ (dual of state constraint)

Step 2: Find optimal risk-adjusted probability by solving the dual:

$$\mu^* = \operatorname{argmax}_{\mu \in \mathcal{M}(p)} \sum_{\omega} \mu_\omega \cdot Q_t(\hat{x}_{t-1}, \omega) - \alpha(p, \mu)$$

For EAVaR with parameters (λ, α) , this is a linear program:

$$\max_{\mu} \sum_{\omega} \mu_\omega \cdot Q_\omega \quad \text{s.t.} \quad \mu_\omega \leq (1 - \lambda)p_\omega + \frac{\lambda p_\omega}{\alpha}, \quad \sum_{\omega} \mu_\omega = 1, \quad \mu \geq 0$$

The solution places maximum weight on the worst (highest-cost) scenarios.

Step 3: Compute risk-averse cut coefficients using the theorem from Section 16.5:

$$\bar{\alpha}_{t-1} = \sum_{\omega \in \Omega_t} \mu_\omega^* \cdot \alpha_t(\omega)$$

$$\bar{\beta}_{t-1} = \sum_{\omega \in \Omega_t} \mu_\omega^* \cdot \beta_t(\omega)$$

Step 4: Add cut to stage $t - 1$:

$$\theta_{t-1} \geq \bar{\alpha}_{t-1} + \bar{\beta}_{t-1}^\top x_{t-1}$$

Note: For pure CVaR ($\lambda = 1$), the optimal μ^* assigns weight only to scenarios with costs at or above VaR_α . For the convex combination ($0 < \lambda < 1$), all scenarios receive some weight.

16.8 Per-Stage Risk Profiles

Risk aversion can vary by stage. The configuration specifies (λ_t, α_t) for each stage:

Scenario	Stage 1-12	Stage 13-60	Stage 61-120
Conservative	$\lambda = 0.5, \alpha = 0.2$	$\lambda = 0.3, \alpha = 0.3$	$\lambda = 0.1, \alpha = 0.5$
Aggressive	$\lambda = 0.1, \alpha = 0.5$	$\lambda = 0.2, \alpha = 0.3$	$\lambda = 0.3, \alpha = 0.2$

16.9 Implementation Notes

Mathematical Concept	Data Model Reference
Risk measure $\rho^{\lambda, \alpha}$	<code>stages.json</code> $\rightarrow$ <code>risk_measure</code> field (per stage)
λ parameter	<code>risk_measure.lambda</code>
α parameter	<code>risk_measure.alpha</code>
Risk-adjusted probabilities $\tilde{p}(\omega)$	Computed at runtime during backward pass
CVaR cut coefficients	Stored in <code>policy/cuts/stage_XXX.bin</code> (same format as risk-neutral)

16.10 Upper Bound with Risk Measures

Important: Monte Carlo simulation cannot directly estimate the upper bound for CVaR problems because: 1. CVaR is computed over the entire distribution, not sample averages 2. The optimal η (VaR threshold) changes with the policy

Solution: Use the inner approximation (SIDP) from Section 8 for true upper bounds with CVaR objectives.

16.11 Reference

Philpott, A.B., de Matos, V.L., & Finardi, E.C. (2013). “On solving multistage stochastic programs with coherent risk measures.” *Operations Research*, 61(4), 957-970. <https://doi.org/10.1287/opre.2013.1200>

16.12 Lower Bound Validity with Risk Measures

Critical Warning: The lower bound computed during SDDP training is **NOT a valid bound** for risk-averse problems.

Why the Lower Bound Fails In risk-neutral SDDP, the lower bound $\underline{z} = V_1(x_0)$ is the optimal value of the first-stage LP, which uses cuts that provide valid outer approximations of the expected future cost. This bound converges to the true optimal expected cost.

In risk-averse SDDP, this property **does not hold** because:

1. **Cuts approximate nested risk measures:** Each cut approximates $\rho_t[V_{t+1}(x_t, \omega)]$, where the risk measure ρ_t depends on the *distribution of costs at that stage*.
2. **The LP optimizes under the wrong distribution:** The first-stage LP optimizes $\rho_1[V_2]$, but the risk-adjusted distribution used in the cuts was computed for the *training* states, not for the optimal first-stage decision.
3. **Nested risk measures are not time-consistent in expectation:** Unlike $\mathbb{E}[\cdot]$, the nested application of CVaR does not satisfy:

$$\rho_1[\rho_2[\dots]] \neq \rho[\text{total cost}]$$

What the “Lower Bound” Represents For risk-averse problems, the value $\underline{z} = V_1(x_0)$ computed by SDDP is:

- A **convergence indicator**: It increases monotonically and plateaus when additional cuts provide no improvement
- **NOT a valid lower bound** on the true risk-averse optimal cost

Recommendations

Purpose	Method
Convergence monitoring	Track z stabilization (bound stalling rule)
Valid lower bound	Use inner approximation (Section 15) evaluated with risk measures
Policy evaluation	Monte Carlo simulation with risk-averse policy decisions

Implementation Note: When risk measures are enabled, convergence reports should label the “lower bound” as “convergence indicator” or explicitly note that it is not a valid bound.

Reference

Shapiro, A. (2011). “Analysis of stochastic dual dynamic programming method.” *European Journal of Operational Research*, 209(1), 63-72.

Part VI: Configuration Reference

17. Configuration-Driven LP Variants

This section provides a comprehensive mapping between POWERS configuration options and their effects on the LP subproblem formulation.

17.1 Block Mode Configuration

Option	Value	LP Effect	Reference
modeling.block_mode	"parallel"	Single water balance per stage, averaged generation	Section 5.1
modeling.block_mode	"chronological"	Per-block storage variables, sequential water balance	Section 5.2

17.2 Hydro Production Function

Option	Value	LP Effect	Reference
modeling.production_function	"constant"	Fixed productivity ρ_h	Section 6.1
modeling.production_function	"linear"	Piecewise-linear head approximation	Section 6.2

17.3 Inflow Non-Negativity Treatment

Option	Value	LP Effect	Reference
modeling.inflow_non_negativity	"none"	No slack, may cause infeasibility	Section 9.2
modeling.inflow_non_negativity	"partial"	Add σ_h^{inf} slack with penalty	Section 9.3

Option	Value	LP Effect	Reference
modeling.inflow_non_negativity_constraint	"true"	Pre-truncate in scenario generation	Section 9.4
modeling.inflow_non_negativity_constraint_with_penalty	"false"	Penalty adjustment slack ξ_h	Section 9.5
modeling.inflow_non_negativity_penalty_cost	float	Penalty coefficient c^{inf} (default: 1000)	Section 9.3

17.4 Cut Management

Option	Value	LP Effect	Reference
training.cut_selection.enabled	bool	Enable/disable cut pruning	Section 11
training.cut_selection.method	"level1"	Keep ever-active cuts	Section 11.3
training.cut_selection.method	"lim1"	Limited memory level-1	Section 11.4
training.cut_selection.method	"domination"	Remove dominated cuts	Section 11.5

17.5 Discount Rate

Option	Location	LP Effect	Reference
transitions[].discount_rate	stages.json	Scale cuts by $\beta_{t \rightarrow t+1}$	Section 13

17.6 Horizon Mode

Option	Value	LP Effect	Reference
horizon.mode	"finite"	Terminal value $V_{T+1} = 0$	Section 2
horizon.mode	"infinite_periodic"	Cycle detection, cut sharing	Section 14
horizon.max_horizon_length	int	Maximum forward pass length	Section 14.7

17.7 Upper Bound Evaluation

Option	Value	LP Effect	Reference
upper_bound_evaluation.enabled	bool	Enable vertex-based inner approx	Section 15
upper_bound_evaluation.method	"lipschitz.mode"	Auto-compute Lipschitz constants	Section 15.4

17.8 Risk Measures

Option	Location	LP Effect	Reference
stages[].risk_measure.type	stages.json	Risk measure selection	Section 16
stages[].risk_measure.lambda	stages.json	Risk aversion weight	Section 16.3
stages[].risk_measure.alpha	stages.json	CVaR confidence level	Section 16.2

17.9 Penalty Coefficients

Option	Default	Objective Term	Reference
modeling.deficit_penalty	10000.0	$c^{def} \cdot \delta_b$	Section 4.1
modeling.spillage_penalty	0.001	$c^{spill} \cdot s_h$	Section 4.5
modeling.inflow_non_negativity_penalty_cost	1000.0	$c^{inf} \cdot \sigma_h^{inf}$	Section 9.3

17.10 Complete Example Configuration

```
{
  "modeling": {
    "block_mode": "chronological",
    "production_function": "fpha",
    "inflow_non_negativity": {
      "method": "penalty",
      "penalty_cost": 1000.0
    },
    "deficit_penalty": 10000.0,
    "spillage_penalty": 0.001
  },
  "training": {
    "forward_passes": 10,
    "cut_selection": {
      "enabled": true,
      "method": "domination",
      "threshold": 0,
      "check_frequency": 10
    },
    "stopping_rules": [
      {"type": "iteration_limit", "limit": 100},
      {"type": "bound_stalling", "iterations": 10, "tolerance": 0.0001}
    ]
  },
  "horizon": {
    "mode": "finite"
  },
  "upper_bound_evaluation": {
    "enabled": true,
    "initial_iteration": 10,
    "interval_iterations": 5
  }
}
```

18. Cross-Reference to Data Model Specification

This section maps each mathematical formulation to the corresponding configuration options and data structures in DATA_MODEL_SPECIFICATION.md.

18.1 Section Mapping

Section	Data Model	Config Path	Data Files
5. Block Formulation	3.2 (Block Mode)	config.json → modeling.block_mode	Stage.blocks[], per-block water balance
6. Production Functions	3.5.1-3.5.4 (Hydro)	hydro_production_models.json, config.json → modeling.production_function	Hydro.productivity, fpha_hyperplanes.parquet, hydro_geometry.parquet
8. PAR(p) Model	3.7 (Inflow Models)	scenarios/inflow_models.parquet	inflowModel struct, Yule-Walker solver
9. Non-Negativity	3.2 (Inflow)	config.json → modeling.inflow_non_negativity	LP slack variables, penalty coefficients
10. Cut Generation	3.5 (Cuts)	N/A (runtime)	Cut struct, dual extraction
11. Cut Selection	3.2 (Cut Selection)	config.json → training.cut_selection	CutPool activity tracking, Cut.active_count
12. Stopping Rules	3.2 (Stopping)	config.json → training.stopping_rules[]	StoppingRule enum, convergence metrics
13. Discount Rate	3.2 (Stages)	stages.json → transitions[].discount_rate	Transition.discount_rate, cut scaling
14. Infinite Horizon	3.2 (Horizon)	config.json → horizon.mode, stages.json cycle	PolicyGraph, cycle cut sharing
15. Inner Approximation	3.2 (Upper Bound)	config.json → upper_bound_evaluation	Vertex struct, Lipschitz constants
16. Risk-Averse CVaR	3.2 (DEFERRED)	stages.json → risk_measure	Risk-adjusted probability computation

18.2 Variable Correspondence

Math Symbol	Field Name	JSON/File Path	Type
v_h	Hydro storage	hydros.json → storage	f64
$\hat{v}_h$	Incoming storage (state)	Internal state vector	Vec<f64>
a_h	Incremental inflow	inflow_models.parquet	f64
$\psi_{m,\ell}$	AR coefficients	inflow_models.parquet → ar_coef_*	[f64; MAX_AR_ORDER]
σ_m	Residual std dev	inflow_models.parquet → residual_std	f64
θ	Future cost variable	LP variable	f64
α_k	Cut intercept	policy/cuts/stage_XXX.bin	f64
β_k	Cut coefficients	policy/cuts/stage_XXX.bin	Vec<f64>
$\beta_{t \rightarrow t+1}$	Discount factor	stages.json → transitions[].discount_rate	f64
L_t	Lipschitz constant	Computed from penalties	f64

18.3 Configuration Quick Reference

Chronological Blocks (Section 5)

```
{
  "modeling": {
    "block_mode": "chronological"
  }
}
```

Effect: Enables per-block storage variables $v_{h,k}$ and sequential water balance constraints.

Discount Rate (Section 13)

```
{
  "stages": [
    {
      "transitions": [
        {"source_id": 0, "target_id": 1, "probability": 1.0, "discount_rate": 0.005}
      ]
    }
  ]
}
```

Effect: Cuts are scaled by $\beta_{t \rightarrow t+1} = 1/(1 + 0.005) \approx 0.995$.

PAR(p) Model (Section 8) Configured via `scenarios/inflow_models.parquet`:

Column	Math Symbol	Description
ar_order	p	Number of AR lags
ar_coef_1 ... ar_coef_12	$\psi_{m,1} \dots \psi_{m,p}$	AR coefficients
seasonal_mean	μ_m	Season mean
residual_std	σ_m	Innovation standard deviation

Inflow Non-Negativity (Section 9)

```
{
  "modeling": {
    "inflow_non_negativity": {
      "method": "penalty",
      "penalty_cost": 1000.0
    }
  }
}
```

Method	Math Formulation	LP Variables
none	Direct AR output	None added
penalty	$a_h + \sigma_h^{inf} = \text{AR output}$	inflow_slack
truncation	$a_h = \max(0, \text{AR output})$	None
truncation_with_penalty	$\eta_h^{adj} = \eta_h + \xi_h$	noise_slack

Cut Selection (Section 11)

```
{
  "training": {
    "cut_selection": {
      "enabled": true,
      "method": "domination",
      "threshold": 0,
      "check_frequency": 10
    }
  }
}
```

Method	Algorithm
level1	Keep cuts active at least once
lm1	Keep most recently active per state
domination	Remove Pareto-dominated cuts

Stopping Rules (Section 12)

```
{
  "training": {
    "stopping_rules": [
      {"type": "iteration_limit", "limit": 50},
      {"type": "simulation", "replications": 100, "period": 20, "distance_tol": 0.01, "bound_tol": 0.0001}
    ],
    "stopping_mode": "any"
  }
}
```

Infinite Horizon (Section 14)

```
{
  "horizon": {
    "mode": "infinite_periodic",
    "max_horizon_length": 240,
    "cycle_discretization_delta": 0.1
  }
}
```

Requirements: - At least one transition must create a cycle - Cycle transitions must have `discount_rate > 0`

Inner Approximation (Section 15)

```
{
  "upper_bound_evaluation": {
    "enabled": true,
    "initial_iteration": 10,
    "interval_iterations": 5,
    "lipschitz": {
      "mode": "auto",
      "fallback_value": 10000.0,
      "scale_factor": 1.1
    }
  }
}
```

18.4 Rust Struct Correspondence

Math Entity	Struct	File Location
Benders cut (α_k, β_k)	Cut	powers-core/src/policy/cut.rs
Cut pool $\mathcal{K}_t$	CutPool	powers-core/src/policy/cut_pool.rs
Vertex $(x^{(i)}, \bar{v}^{(i)})$	Vertex	powers-core/src/policy/vertex.rs
Stage subproblem	Subproblem	powers-core/src/lp/subproblem.rs
State vector x_t	StateVector	powers-core/src/state.rs
PAR(p) model	InflowModel	powers-core/src/stochastic/inflow.rs
Risk measure $\rho^{\lambda, \alpha}$	RiskMeasure	powers-core/src/risk.rs

Math Entity	Struct	File Location
Transition graph	PolicyGraph	powers-core/src/graph.rs

Summary

This document provides the complete mathematical foundations for the POWE.RS SDDP solver:

Part I: SDDP Algorithm Foundation 1. **Introduction:** Document purpose, notation conventions, problem context 2. **SDDP Algorithm Overview:** Forward/backward passes, convergence, policy graphs

Part II: Stage Subproblem Formulation 3. **Notation and Sets:** Index sets, parameters, variables, duals 4. **Base LP Formulation:** Objective, load balance, water balance, AR dynamics, generation constraints 5. **Block Formulation Variants:** Parallel blocks (default), chronological blocks 6. **Hydro Production Function Models:** Constant productivity, FPHA 7. **Equipment-Specific Formulations:** Thermal plants, transmission lines, contracts, pumping stations

Part III: Stochastic Modeling 8. **PAR(p) Inflow Model:** Yule-Walker equations, seasonal correlations, model fitting 9. **Inflow Non-Negativity Methods:** None, penalty, truncation, truncation with penalty

Part IV: Cut Management and Convergence 10. **Cut Generation and Aggregation:** Dual extraction, coefficient computation, single-cut aggregation 11. **Cut Selection Strategies:** Level-1, LML1, domination-based 12. **Stopping Rules:** Iteration limit, time limit, statistical, bound stalling, simulation-based

Part V: Advanced Formulations 13. **Discount Rate:** Discounted Bellman equation, cumulative discounting, cut scaling 14. **Infinite Periodic Horizon:** Cycle detection, cut sharing, fixed-point iteration 15. **Upper Bound Evaluation LP:** Vertex-based inner approximation, Lipschitz interpolation 16. **Risk-Averse SDDP (CVaR):** Coherent risk measures, probability reweighting, dual representation

Part VI: Configuration Reference 17. **Configuration-Driven LP Variants:** Complete mapping of config options to LP changes 18. **Cross-Reference to Data Model:** Variable correspondence, Rust struct mapping

Appendices - Appendix A: Notation reference - **Appendix B:** SDDP algorithm pseudocode - **Appendix C:** Deferred features (GNL, batteries, multi-cut, Markovian)

Each formulation follows SDDP.jl notation conventions and is designed for production-scale hydrothermal dispatch problems.

Appendix A: Notation Reference

A.1 Index Sets

Symbol	Description	Typical Size
$\mathcal{T}$ or $t \in \{1, \dots, T\}$	Stages	60-120
$\mathcal{K}$ or k	Blocks within stage	1-24
$\mathcal{H}$	Hydro plants	160
$\mathcal{B}$	Buses	4-10
Ω_t	Scenario realizations at stage t	10-50
$\mathcal{K}_t$	Cuts at stage t	100-10,000
$\mathcal{V}_t$	Vertices at stage t	100-10,000

A.2 State Variables

Symbol	Units	Bounds	Description
v_h	hm ³	$[\underline{V}_h, \bar{V}_h]$	Reservoir storage
$a_{h,\ell}$	m ³ /s	free	AR lag ℓ inflow
θ	\$	≥ 0	Future cost (epigraph)

A.3 Control Variables

Symbol	Units	Bounds	Description
$q_{h,k}$	m ³ /s	$[\underline{Q}_h, \bar{Q}_h]$	Turbined flow
$s_{h,k}$	m ³ /s	≥ 0	Spillage
$u_{h,k}$	m ³ /s	$[0, \bar{U}_h]$	Diversion/bypass flow
$r_{h,k}$	m ³ /s	≥ 0	Water withdrawal (consumptive)
$o_{h,k}$	m ³ /s	-	Total downstream outflow: $o = q + s$
$g_{h,k}$	MW	$[\underline{G}_h, \bar{G}_h]$	Hydro generation
$g_{j,k,s}$	MW	$[\underline{G}_j, \bar{G}_j]$	Thermal generation (plant j , segment s)
$\chi_{c,k}^{in}$	MW	$[0, \bar{C}_c]$	Contract import
$\chi_{c,k}^{out}$	MW	$[0, \bar{C}_c]$	Contract export
$\delta_{b,k}$	MW	≥ 0	Deficit (load shedding)
μ_ω	-	$[0, 1]$	Risk-adjusted probability (CVaR)

A.4 Parameters

Symbol	Units	Source	Description
μ_m	m ³ /s	<code>inflow_models.parquet</code>	Seasonal mean inflow
$\psi_{m,\ell}$	-	<code>inflow_models.parquet</code>	AR coefficient
σ_m	m ³ /s	<code>inflow_models.parquet</code>	Residual std dev
ρ_h	MW/(m ³ /s)	<code>hydros.json</code>	Productivity
$\beta_{t \rightarrow t+1}$	-	<code>stages.json</code>	Discount factor
L_t	\$/hm ³	Computed	Lipschitz constant

A.5 Dual Variables (Cut Coefficients)

Symbol	Constraint	Cut Coefficient
π_h^{wb}	Water balance	$\beta_h^v = \zeta \cdot \pi_h^{wb}$
$\pi_{h,\ell}^{lag}$	AR lag fixing	$\beta_{h,\ell}^{lag} = \pi_{h,\ell}^{lag}$
λ_k	Cut constraint	Activity indicator

Appendix B: SDDP Algorithm Pseudocode

B.1 Main Training Loop

Algorithm: SDDP Training

- **Input:** Policy graph G , initial state x_0 , stopping rules R

- **Output:** Trained policy (cuts K , vertices V)

Initialize:

- $K_t = \emptyset$ for all stages t (cut pools)
- $V_t = \emptyset$ for all stages t (vertex pools, if inner approximation enabled)
- $k = 0$ (iteration counter)

Repeat:

1. Increment: $k = k + 1$
2. **Forward Pass:** For $m = 1$ to M forward passes:
 - Set $x = x_0$, initialize **trajectory** = []
 - For $t = 1$ to T :
 - Sample $\omega = \text{sample_scenario}(t)$
 - Solve $(x', \text{cost}) = \text{solve_subproblem}(t, x, \omega, K_t)$
 - Record **trajectory.append**(($t, x, x', \omega, \text{cost}$)), update $x = x'$
 - Store trajectory in **trial_points**
3. **Backward Pass:** For $t = T - 1$ down to 1:
 - For each $(t, \hat{x}, _, _, _)$ in **trial_points**:
 - For each $\omega \in \Omega_t$:
 - * Solve $(Q, \pi) = \text{solve_subproblem_with_duals}(t + 1, \hat{x}, \omega, K_{t+1})$
 - * Compute cut $(\alpha, \beta) = \text{compute_cut_coefficients}(Q, \pi, \hat{x})$
 - * Store in **scenario_cuts**
 - Aggregate $(\bar{\alpha}, \bar{\beta}) = \text{aggregate_cuts}(\text{scenario_cuts}, \text{risk_measure}_t)$
 - Add to **cuts_for_stage**
 - Update $K_t = K_t \cup \text{cuts_for_stage}$
4. **Convergence Check:**
 - Compute $\text{lower_bound} = \text{solve_first_stage_LP}(x_0, K_1)$
 - If needed: $\text{upper_bound} = \text{simulate_policy}(x_0, K, \text{num_simulations})$
 - Compute $\text{gap} = (\text{upper_bound} - \text{lower_bound}) / |\text{upper_bound}|$
 - If enabled: $K = \text{select_active_cuts}(K)$

Until any stopping rule is satisfied: **any**(**rule.evaluate**($k, \text{lower_bound}, \text{upper_bound}$) for **rule** in R)

Return K, V

B.2 Subproblem Solve

Function: **solve_subproblem**($t, x_{\text{prev}}, \omega, \text{cuts}$)

- **Input:** Stage t , incoming state x_{prev} , scenario ω , cuts K_t
- **Output:** (outgoing state x , stage cost c)

Build LP:

Variables:

- x_t = state variables (storage, AR lags)
- u_t = control variables (generation, flow, deficit)
- θ_t = future cost variable

Objective:

$$\min c_t^\top u_t + \theta_t$$

Constraints:

- Load balance: $\sum \text{generation} = \text{demand}_t(\omega)$
- Water balance: $x_t.\text{storage} = x_{\text{prev}}.\text{storage} + \text{inflow}_t(\omega) - \text{outflow}$
- AR dynamics: $x_t.\text{lags} = \text{update_lags}(x_{\text{prev}}.\text{lags}, \text{inflow}_t(\omega))$

- Bounds: $x_t \in X_t, u_t \in U_t$
- Benders cuts: For each (α_k, β_k) in cuts:

$$\theta_t \geq \beta_{t \rightarrow t+1} \cdot (\alpha_k + \beta_k^\top x_t)$$

Solve LP with HiGHS and Return $(x_t^*, c_t^\top u_t^* + \theta_t^*)$

B.3 Cut Coefficient Computation

Function: `compute_cut_coefficients`($Q, \pi, \hat{x}$)

- **Input:** Optimal objective Q , dual multipliers π , trial point $\hat{x}$
- **Output:** Cut coefficients (α, β)

1. **Storage coefficients** (from water balance duals):

- For each hydro h :

$$\beta_h^v = \pi_h^{wb} \cdot \zeta$$

where ζ is the time conversion factor

2. **AR lag coefficients** (from lag fixing duals):

- For each hydro h , lag ℓ :

$$\beta_{h,\ell}^{lag} = \pi_{h,\ell}^{lag}$$

3. **Intercept** (makes cut pass through trial point):

$$\alpha = Q - (\beta^v)^\top \hat{x}.\text{storage} - (\beta^{lag})^\top \hat{x}.\text{lags}$$

4. **Return** $(\alpha, \beta^v, \beta^{lag})$

Appendix C: Deferred Features

This appendix documents features that are planned but not yet implemented in the POWERS SDDP solver.

C.1 GNL Thermal Plants

Status: DEFERRED

Description: Gas Natural Liquefeito (GNL) thermal plants have complex operational constraints including: - Minimum take-or-pay contracts - Variable fuel costs based on LNG spot market - Start-up and shutdown constraints - Fuel inventory management

Planned Formulation: - Binary variables for unit commitment (requires MIP solver integration) - Fuel inventory balance constraints - Contract fulfillment constraints - Piecewise-linear fuel cost functions

Reference: CEPEL NEWAVE/DECOMP GNL modeling documentation.

C.2 Battery Energy Storage Systems

Status: DEFERRED

Description: Grid-scale batteries with: - State-of-charge management - Charge/discharge efficiency losses - Degradation modeling (cycle counting) - Capacity fade over time

Planned Formulation:

Variables: | Variable | Domain | Units | Description | | $e_{b,k}$ | $[0, \bar{E}_b]$ | MWh | State of charge | | $p_{b,k}^{ch}$ | ≥ 0 | MW | Charging power | | $p_{b,k}^{dis}$ | ≥ 0 | MW | Discharging power |

Energy Balance:

$$e_{b,k} = e_{b,k-1} + \eta^{ch} \cdot p_{b,k}^{ch} \cdot \Delta t - \frac{p_{b,k}^{dis} \cdot \Delta t}{\eta^{dis}}$$

Load Balance Contribution:

$$\sum_{b \in \mathcal{B}} (p_{b,k}^{dis} - p_{b,k}^{ch}) \text{ added to generation}$$

C.3 Multi-Cut Formulation

Status: DEFERRED

Description: Alternative to single-cut aggregation that creates one cut per scenario.

Formulation:

Instead of a single aggregated future cost variable θ , introduce per-scenario variables θ_ω :

$$\theta = \sum_{\omega \in \Omega_t} p(\omega) \cdot \theta_\omega$$

With per-scenario cuts:

$$\theta_\omega \geq \alpha_k(\omega) + \beta_k(\omega)^\top x \quad \forall k, \omega$$

Trade-offs: - **Advantages:** Potentially faster early convergence, better for problems with few scenarios - **Disadvantages:** More constraints per iteration, higher LP solve time, more memory

Reference: Birge, J.R. (1985). “Decomposition and partitioning methods for multistage stochastic linear programs.” *Operations Research*, 33(5), 989-1007.

C.4 Markovian Policy Graphs

Status: DEFERRED

Description: Extension to handle scenario-dependent transitions (e.g., different inflow regimes).

Current Limitation: POWERS assumes stage-wise independent scenarios. The same scenario tree structure is used regardless of which scenario was realized in the previous stage.

Planned Extension: - Markov chain over “regimes” (e.g., wet/dry/normal) - Transition probabilities between regimes - Regime-dependent inflow distributions - Cut sharing across nodes in same regime

Formulation:

Let M be a Markov chain with states $\mathcal{M} = \{1, \dots, m\}$ and transition matrix P .

The policy graph becomes: - Nodes: (t, r) for stage t and regime r - Edges: $(t, r) \rightarrow (t+1, r')$ with probability $P_{r,r'}$

Value function approximation:

$$V_{t,r}(x) \approx \max_{k \in \mathcal{K}_{t,r}} \{\alpha_k + \beta_k^\top x\}$$

Cuts are regime-specific and only shared within the same regime.

Reference: Philpott, A.B., & de Matos, V.L. (2012). “Dynamic sampling algorithms for multi-stage stochastic programs with risk aversion.” *European Journal of Operational Research*, 218(2), 470-483.

C.5 Non-Controllable Sources (Wind/Solar)

Status: DEFERRED

Description: Stochastic renewable generation with: - Availability factors correlated with inflows - Curtailment decisions - Capacity credit calculations

Planned Formulation:

Variables: | Variable | Domain | Units | Description | |-----|-----|-----|-----| | $g_{r,k}^{nc}$ | $[0, \bar{G}_r \cdot \alpha_r(\omega)]$ | MW | Non-controllable generation | | $\kappa_{r,k}$ | ≥ 0 | MW | Curtailment |

Generation Constraint:

$$g_{r,k}^{nc} + \kappa_{r,k} = \bar{G}_r \cdot \alpha_r(\omega)$$

where $\alpha_r(\omega) \in [0, 1]$ is the stochastic availability factor.

Curtailment Penalty:

$$+ \sum_{r \in \mathcal{R}} c^{curt} \cdot \kappa_{r,k} \cdot \Delta t_k$$

C.6 FPFA Enhancements

Status: DEFERRED (Partial - core FPFA implemented)

Description: Advanced extensions to the FPFA (Four-Point Head Approximation) model for improved accuracy in hydroelectric production function modeling.

C.6.1 Variable Efficiency Curves (q) **Current:** Constant turbine-generator efficiency η_{ref} .

Enhancement: Flow-dependent efficiency using characteristic curves:

$$\eta(q) = \eta_{max} \times f\left(\frac{q}{q_{nom}}\right)$$

where f is a hill chart approximation, typically:

$$f(x) = a_0 + a_1x + a_2x^2 + a_3x^3 \quad \text{for } x \in [x_{min}, 1]$$

Data Requirements: - Efficiency curve coefficients in `hydro_production_data.parquet` - `efficiency_type = "flow_dependent"` - `efficiency_coeffs = [a_0, a_1, a_2, a_3]`

Impact on FPFA: - Hyperplane fitting must use $\eta(q_i)$ at each grid point - Increases nonlinearity captured by the approximation - More planes may be needed for same accuracy

C.6.2 Pumped Hydro Production Function **Current:** Pumping modeled separately from generation (see pumping stations).

Enhancement: Unified production function for reversible hydro plants:

Generation Mode (standard FPFA):

$$g_{h,k}^{gen} \leq \phi(v, q^{gen}, s)$$

Pumping Mode:

$$p_{h,k}^{pump} \geq \phi_{pump}(v, q^{pump}, s)$$

Note the reversed inequality - pumping power consumption *increases* with head.

Pumping Production Function:

$$p_{pump} = \frac{\rho \times q^{pump} \times h_{net,pump}}{\eta_{pump}}$$

where: - $h_{net,pump}$ = pumping head (downstream $\rightarrow$ upstream) - η_{pump} = pumping efficiency (typically 0.85-0.90, lower than generation)

Hyperplane Form:

$$p_{h,k}^{pump} \geq \gamma_{0,pump}^m + \gamma_{v,pump}^m \cdot v_h^{avg} + \gamma_{q,pump}^m \cdot q_{h,k}^{pump}$$

Operational Constraints: - Mutual exclusion: $q_{h,k}^{gen} \times q_{h,k}^{pump} = 0$ (nonlinear) - Alternative: Big-M or SOS1 constraints (introduces integer variables) - POWE.RS approach: Allow simultaneous gen/pump with high penalty (relaxation)

C.6.3 Dynamic FPFA Recomputation **Current:** FPFA hyperplanes fixed per stage configuration.

Enhancement: Recompute hyperplanes based on expected operating region.

Motivation: The convex hull approximation quality depends on how well the fitting grid matches actual operation. If a hydro operates far from the fitting window center, the approximation may be poor.

Approach: 1. **Training phase:** Track volume distribution per stage across forward passes 2. **Update fitting windows:** Adjust $[v_{min}, v_{max}]$ to cover observed operation 3. **Refit hyperplanes:** Generate new planes for updated windows 4. **Cut validity concern:** Changing hyperplanes may invalidate existing cuts

Mitigation Strategies: - Only narrow fitting windows (ensures new FPFA is more restrictive) - Rebuild cuts after significant FPFA updates - Use conservative (wider) windows during training

Configuration (future):

```
{
  "fpha_config": {
    "dynamic_recomputation": {
      "enabled": true,
      "recompute_every_n_iterations": 50,
      "window_adaptation": "narrow_only",
      "percentile_margin": 5
    }
  }
}
```

Warning: Dynamic recomputation changes the LP structure across iterations. This may affect SDDP convergence guarantees. Use with caution and validate results against fixed-FPFA runs.

C.7 Temporal Scope Decoupling

Status: DEFERRED

Description: Advanced temporal decomposition feature inspired by SPARHTACUS that decouples the physical time resolution (decision dynamics) from the SDDP stage decomposition (Benders cut generation points). This enables flexible multi-resolution modeling with controlled cut growth.

C.7.1 Motivation Problem with Standard SDDP Temporal Resolution:

In conventional SDDP implementations (including POWERS current design), the following temporal scopes are tightly coupled:


1. **Stage** = SDDP decomposition unit (where state variables live and cuts are generated)
2. **Decision period** = Physical time resolution for operational constraints
3. **Stochastic process** = Uncertainty realization base

This coupling creates a fundamental trade-off: - **Fine temporal resolution** → Accurate physics modeling BUT exponential cut growth - **Coarse temporal resolution** → Manageable cuts BUT poor short-term dynamics

Example Dilemma: - For a 5-year horizon: Monthly stages → 60 stages, manageable cuts - But monthly resolution cannot capture weekly cycling patterns in the first month - Switching to weekly resolution → 260+ stages, cut explosion

Temporal Scope Decoupling Solution:

Allow stages to contain **multiple internal decision periods** with full temporal dynamics, while only generating Benders cuts at stage boundaries:



Benefits: - Week 1 can be decomposed into weekly decisions within monthly stage 0 - Benders cuts only at stage boundaries (60 cuts for 5 years) - Intra-stage dynamics captured without state dimension explosion - Enables configurations like: “First month weekly, rest monthly, year 2+ annually”

C.7.2 Three Independent Temporal Scopes (SPARHTACUS Concept) Following SPARHTACUS nomenclature ([Norus SPTcpp](#)):

Scope	Portuguese Term	POWE.RS Current	Purpose
Optimization Period	Período de otimização	stages[t]	SDDP decomposition unit, Benders cut generation
Study Period	Período de estudo	<i>(coupled to stage)</i>	Physical time resolution for constraints and decisions
Stochastic Process Period	Período do processo estocástico	inflow_models.parquet per stage	Base for uncertainty realization

Key Insight: These three can be **independent**:

Example: Medium-term planning (DECOMP-like)

```
{
  "stages": [
    {
      "id": 0,
      "optimization_period": {"start": "2024-01-01", "duration": "1M"},
      "study_periods": [
        {"id": 0, "duration": "7d", "blocks": [...]}, // Week 1
        {"id": 1, "duration": "7d", "blocks": [...]}, // Week 2
        {"id": 2, "duration": "7d", "blocks": [...]}, // Week 3
        {"id": 3, "duration": "10d", "blocks": [...]} // Remaining days
      ],
      "stochastic_periods": [
        {"id": 0, "duration": "14d"}, // First two weeks
        {"id": 1, "duration": "16d"} // Remaining days
      ]
    },
    {
      "id": 1,
      "optimization_period": {"start": "2024-02-01", "duration": "1M"},
      "study_periods": [
        {"id": 0, "duration": "1M", "blocks": [...]} // Single monthly period
      ],
      "stochastic_periods": [
        {"id": 0, "duration": "1M"}
      ]
    }
  ]
}
```

Characteristics: - Stage 0: 4 study periods (weekly detail), 2 stochastic periods (biweekly inflows) - Stage 1: 1 study period (monthly), 1 stochastic period (monthly inflows) - Benders cuts: Only 2 cuts (at stage boundaries), not 5

C.7.3 Mathematical Formulation Impact **Current POWERS Stage Subproblem** (single decision period per stage):

$$V_t(x_t, \omega_t) = \min_{y_t} \left\{ c_t^\top y_t + \mathbb{E}_{\omega_{t+1}} [V_{t+1}(x_{t+1}, \omega_{t+1})] \right\}$$

Subject to: - $A_t y_t = b_t(\omega_t) - B_t x_t - y_t \geq 0$

Extended Multi-Period Stage Subproblem:

Let stage t contain K_t decision periods indexed by $k = 1, \dots, K_t$.

$$V_t(x_t, \{\omega_{t,k}\}_{k=1}^{K_t}) = \min_{\{y_{t,k}\}_{k=1}^{K_t}} \left\{ \sum_{k=1}^{K_t} c_{t,k}^\top y_{t,k} + \mathbb{E}[V_{t+1}(x_{t+1}, \cdot)] \right\}$$

Subject to: - **Period 1 constraints:** $A_{t,1} y_{t,1} = b_{t,1}(\omega_{t,1}) - B_t x_t$ - **Period k constraints:** $A_{t,k} y_{t,k} + D_{t,k} y_{t,k-1} = b_{t,k}(\omega_{t,k})$ for $k = 2, \dots, K_t$ - **State transition:** $x_{t+1} = E_t y_{t,K_t} + F_t x_t - y_{t,k} \geq 0$ for all k

Key Differences: - **LP variables:** Now include all periods within the stage ($N_{vars} \times K_t$) - **State variable:** Still only x_{t+1} at stage boundary (dimension unchanged!) - **Cuts:** Generated only with respect to x_t (start of stage), not intermediate periods - **Coupling matrix $D_{t,k}$:** Links decisions across periods (e.g., reservoir continuity)

Hydro Water Balance Example (parallel blocks $\rightarrow$ chronological periods):

Period 1 (receives incoming state):

$$v_{h,1} = \hat{v}_h + \zeta_1 [a_{h,1}(\omega_{t,1}) + \text{net_flows}_{h,1}]$$

Period k (sequential continuity):

$$v_{h,k} = v_{h,k-1} + \zeta_k [a_{h,k}(\omega_{t,k}) + \text{net_flows}_{h,k}] \quad k = 2, \dots, K_t$$

State transition (to next stage):

$$v_h^{next} = v_{h,K_t}$$

Benders Cut (end of period K_t):

$$\theta_t \geq \alpha_i + \beta_i^\top \hat{v}_h \quad (\text{dual extracted from period 1 constraint})$$

C.7.4 Data Model Changes Required **Current:** `stages.json` with single start/end date and blocks array

Extended: Add `periods` array within each stage

```
{
  "stages": [
    {
      "id": 0,
      "optimization_period": {
        "start_date": "2024-01-01",
        "end_date": "2024-02-01"
      },
      "periods": [
        {
          "id": 0,
          "duration_days": 7,
          "blocks": [
            {"id": 0, "name": "LEVE", "hours": 56},
            {"id": 1, "name": "MEDIA", "hours": 56},
            {"id": 2, "name": "PESADA", "hours": 56}
          ]
        },
        {
          "id": 1,
          "duration_days": 7,
          "blocks": [...]
        },
        {
          "id": 2,
          "duration_days": 7,
          "blocks": [...]
        },
        {
          "id": 3,
          "duration_days": 10,
          "blocks": [...]
        }
      ],
      "stochastic_process": {
        "realization_periods": [
          {"id": 0, "duration_days": 14}, // Periods 0-1 share same realization
          {"id": 1, "duration_days": 17} // Periods 2-3 share same realization
        ]
      }
    }
  ]
}
```

```

    },
    "num_scenarios": 20
  }
]
}

```

Validation Rules: 1. Sum of `periods[].duration_days` must equal stage duration 2. Stochastic process periods must align with study period boundaries 3. Blocks now defined per period, not per stage 4. Each period can have different block structure

Scenario Generation Changes: - Stochastic process sampled at `realization_periods` boundaries - Multiple study periods can share the same stochastic realization - Inflow models now indexed by (`stage_id`, `stochastic_period_id`)

C.7.5 LP Subproblem Size Impact **Current POWERS** (parallel blocks, 3 blocks per stage):

Component	Count	Formula
Hydro storage vars	N_h	160
Hydro flow vars	$3N_h$	480 (per block)
Total stage vars	~ 1500	$3N_h + 3N_{th} + \dots$
Water balance constraints	N_h	160 (averaged)

Extended Multi-Period (4 periods, 3 blocks each):

Component	Count	Formula
Hydro storage vars	$(K_t - 1)N_h + N_h$	640 (inter-period) + 160 (final state)
Hydro flow vars	$3K_t N_h$	1920
Total stage vars	~ 6000	$4 \times$ current
Water balance constraints	$K_t N_h$	640 (sequential)

Trade-off: - **LP size:** Increases by factor of K_t (number of periods) - **Cut count:** Remains constant (only at stage boundaries) - **Net effect:** For long horizons, LP solve time increase is dominated by cut reduction

Example Scaling (5-year horizon):

Configuration	Stages	Avg Periods/Stage	Total LPs	Avg LP Size	Cut Pool
Standard monthly	60	1	60	1500 vars	~ 1200 cuts @ 20 iter
Hybrid (4 weekly, rest monthly)	60	1.05	60	~ 1600 vars	~ 1200 cuts
All weekly	260	1	260	1500 vars	~ 5200 cuts @ 20 iter

The hybrid approach achieves weekly detail in critical periods without cut explosion.

C.7.6 Implementation Considerations **Algorithm Changes:** 1. **Forward Pass:** Simulate through all periods within each stage sequentially 2. **Backward Pass:** Extract duals from period 1's coupling constraints 3. **Cut Generation:** Cuts reference stage-boundary state only 4. **State Dimension:** Unchanged (no new state variables for intermediate periods)

SDDP Convergence Properties: - Preserved: Multi-period formulation is still a valid SDDP decomposition - The “stage” is now a larger LP, but decomposition structure remains - Convergence guarantees from standard SDDP theory apply

Cut Sharing and Aggregation: - Compatible with existing cut aggregation (Section 10.4) - Cut dimensionality unchanged (still based on x_t at stage start) - Can apply level-1 cuts across stages as before

Compatibility with Chronological Blocks (Section 5.2): - Periods provide **inter-period** dynamics - Blocks within each period can still be parallel or chronological - Nested temporal structure: Stage $\rightarrow$ Periods $\rightarrow$ Blocks

Memory and Solve Time: - LP matrix size: $O(K_t)$ increase per stage - Subproblem solve: $O(K_t^{1.5})$ to $O(K_t^2)$ depending on sparsity - For moderate $K_t \leq 10$: Still faster than adding equivalent stages

C.7.7 Use Cases and Configuration Examples Case 1: Short-term Planning with Daily Detail

```
{
  "stages": [
    {
      "id": 0,
      "comment": "First week with daily resolution",
      "periods": [{"id": k, "duration_days": 1} for k in 0..6],
      "stochastic_process": {"realization_periods": [{"id": 0, "duration_days": 7}]}
    },
    {
      "id": 1,
      "comment": "Rest of month with weekly resolution",
      "periods": [{"id": k, "duration_days": 7} for k in 0..3]
    }
  ]
}
```

Case 2: Medium-term Planning (DECOMP-like)

Month 1: 4 weekly periods in stage 0
 Months 2-12: 1 monthly period per stage (11 stages)
 Years 2-5: 1 annual period per stage (4 stages)
 Total: 16 stages (vs. 260 if all weekly)

****Case 3: Stochastic Process Resolution** Ad'json { "stages": [{ "comment": "Detailed decisions, coarse stochastic process", "periods": [{ "id": 0, "duration_days": 1 }, { "id": 1, "duration_days": 1 }, ...], "stochastic_process": { "realization_periods": [{ "id": 0, "duration_days": 7 } // Single weekly inflow for all 7 daily periods] } }] } "

This enables daily operational constraints with weekly inflow uncertainty.

C.7.8 Comparison to Existing Chronological Blocks **Current Chronological Blocks** (Section 5.2): - Intra-stage storage dynamics between blocks - All blocks share same stochastic realization - Blocks typically parallel in time (peak/off-peak patamares) - Limited to ~3-24 blocks before LP becomes unwieldy

Proposed Multi-Period Stages: - Inter-period storage dynamics (strictly sequential) - Periods can have independent stochastic realizations - Periods always chronological (time-ordered) - Designed for 4-10 periods per stage

Relationship: - Periods Blocks: Each period contains blocks - Orthogonal features that can be combined: - Parallel blocks within sequential periods - Chronological blocks within sequential periods

C.7.9 Open Questions and Design Decisions

1. **Period Duration Constraints:** Should all periods within a stage have equal duration? Or allow variable (as shown in examples)?
 - **Proposal:** Allow variable for flexibility
2. **Stochastic Process Independence:** Should we enforce that stochastic realization periods align exactly with study period boundaries?

- **Proposal:** Yes, to avoid ambiguity in uncertainty propagation
3. **Backward Pass Dual Extraction:** Which period's constraints provide duals for cuts?
 - **Answer:** Period 1, as it couples to incoming state $\hat{x}_t$
 4. **Initial Conditions:** How to specify inflow history when stages have multiple periods?
 - **Proposal:** Pre-study stages remain single-period for simplicity
 5. **Markovian Transitions** (Section C.4): How do multi-period stages interact with Markov chains?
 - **Proposal:** Markov state transitions only at stage boundaries, not between periods

C.7.10 References and Related Work

- **SPARHTACUS/SPTcpp:** [Escopo Temporal](#) - Norus Tecnologia's implementation
- **SDDP.jl:** Supports flexible graph structures but not explicit temporal scopes
- **DECOMP:** Brazilian official model uses weekly/monthly hybrid decomposition
- **Pereira, M.V.F., & Pinto, L.M.V.G. (1991):** "Multi-stage stochastic optimization applied to energy planning" - Original SDDP paper with monthly stages

End of Document